



Telecom Insights

Predicting Churn, Preserving Value

Advanced Machine Learning and
Explainable AI for Telecom Retention
Strategies

PRESENTED BY

Audrey Chang

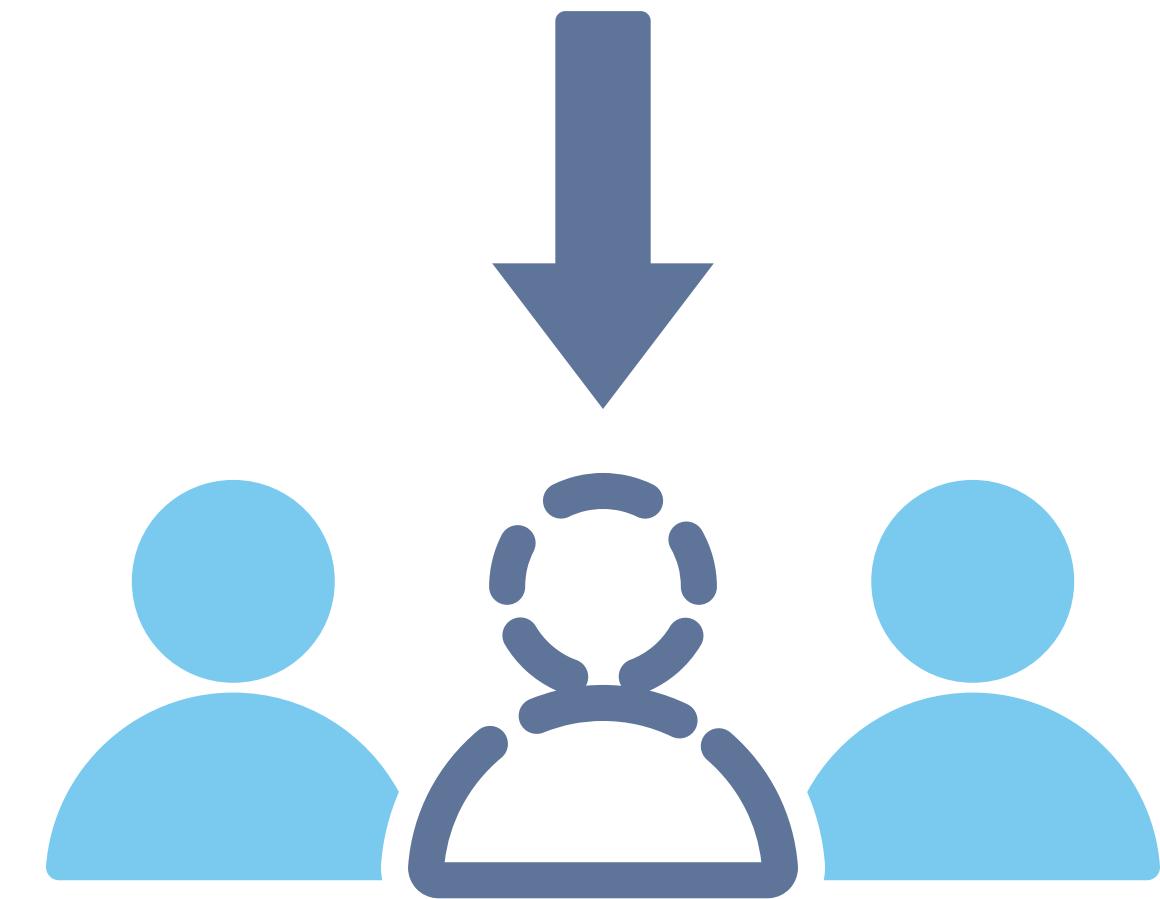


Table of Contents

1. Motivation

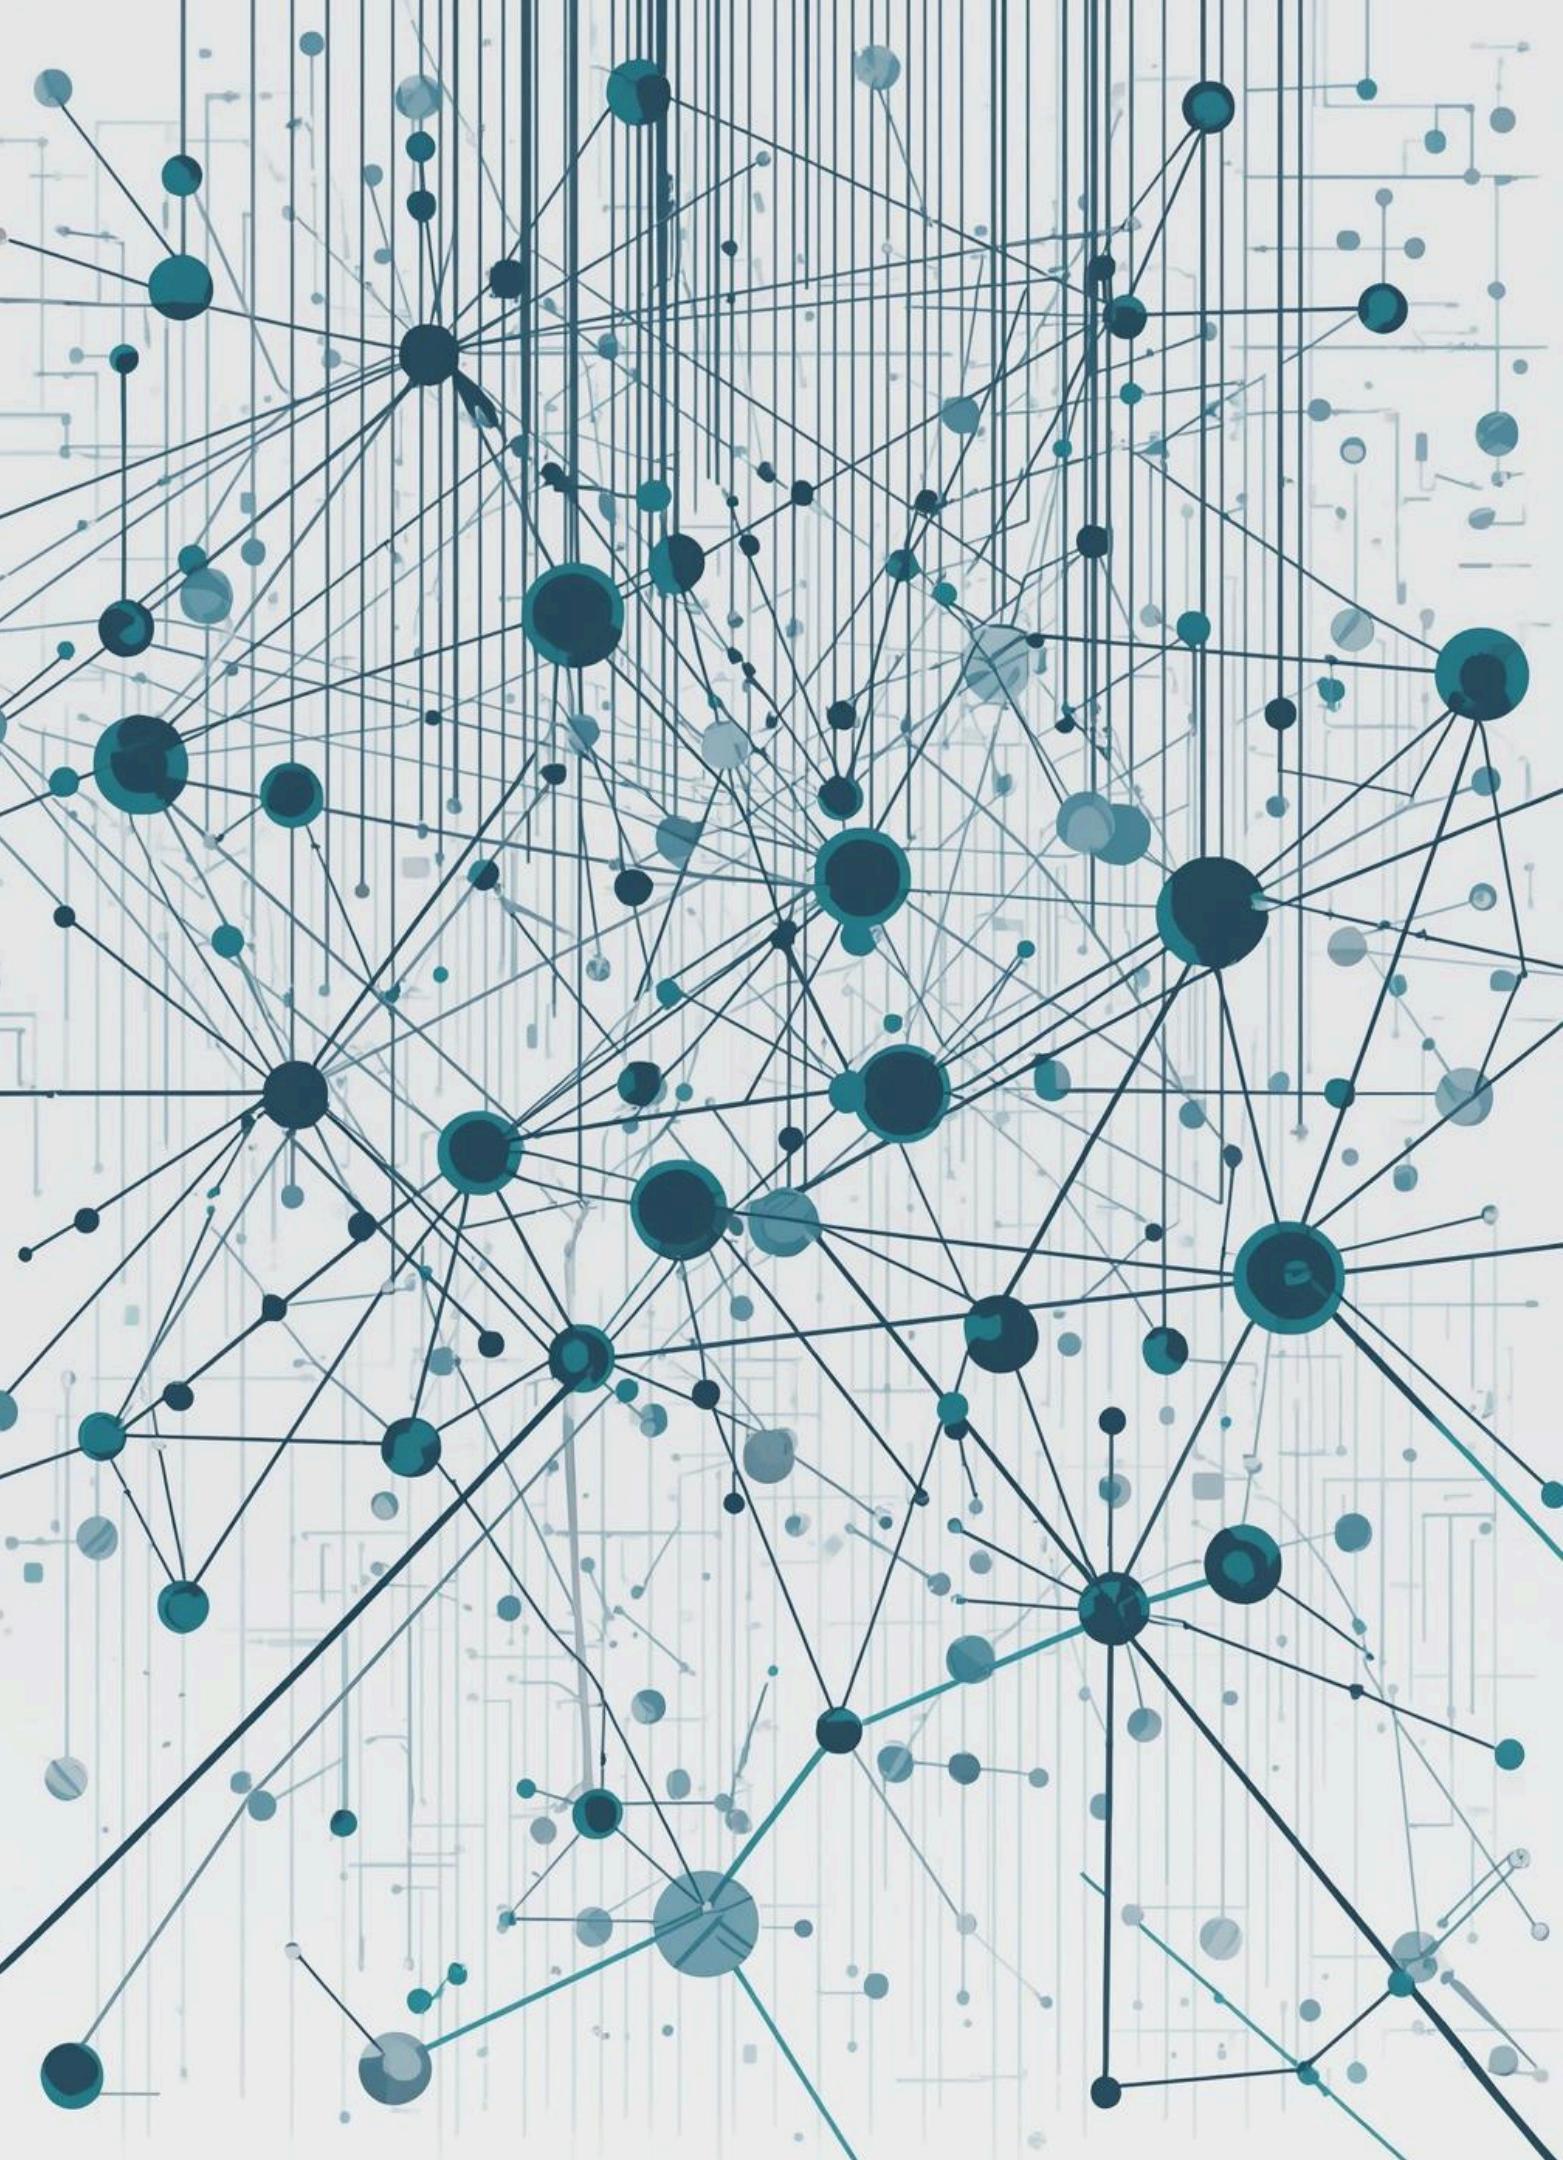
2. Setup

3. Models

4. Results

5. Insights

6. Wrap-up



Problem and Research Questions

- Problem: Telecom customer churn is costly; retention budgets are limited → need accurate and actionable risk ranking.
- Goal: Build an end-to-end churn prediction pipeline under realistic constraints (class imbalance, leakage control, reproducibility)

ENGINEERING RIGOR

Leakage-controlled preprocessing, consistent train/val/test splits, and reproducible artifact logging (configs, models, metrics)

MODEL COMPARISON

Benchmark a calibrated GLM baseline against tree/GBM models, deep tabular networks, and a stacking ensemble

EXPLAINABILITY IMPORTANCE

Translate interpretable drivers and failure modes (SHAP + error analysis) into retention playbooks and budget-aware targeting



Data & Evaluation

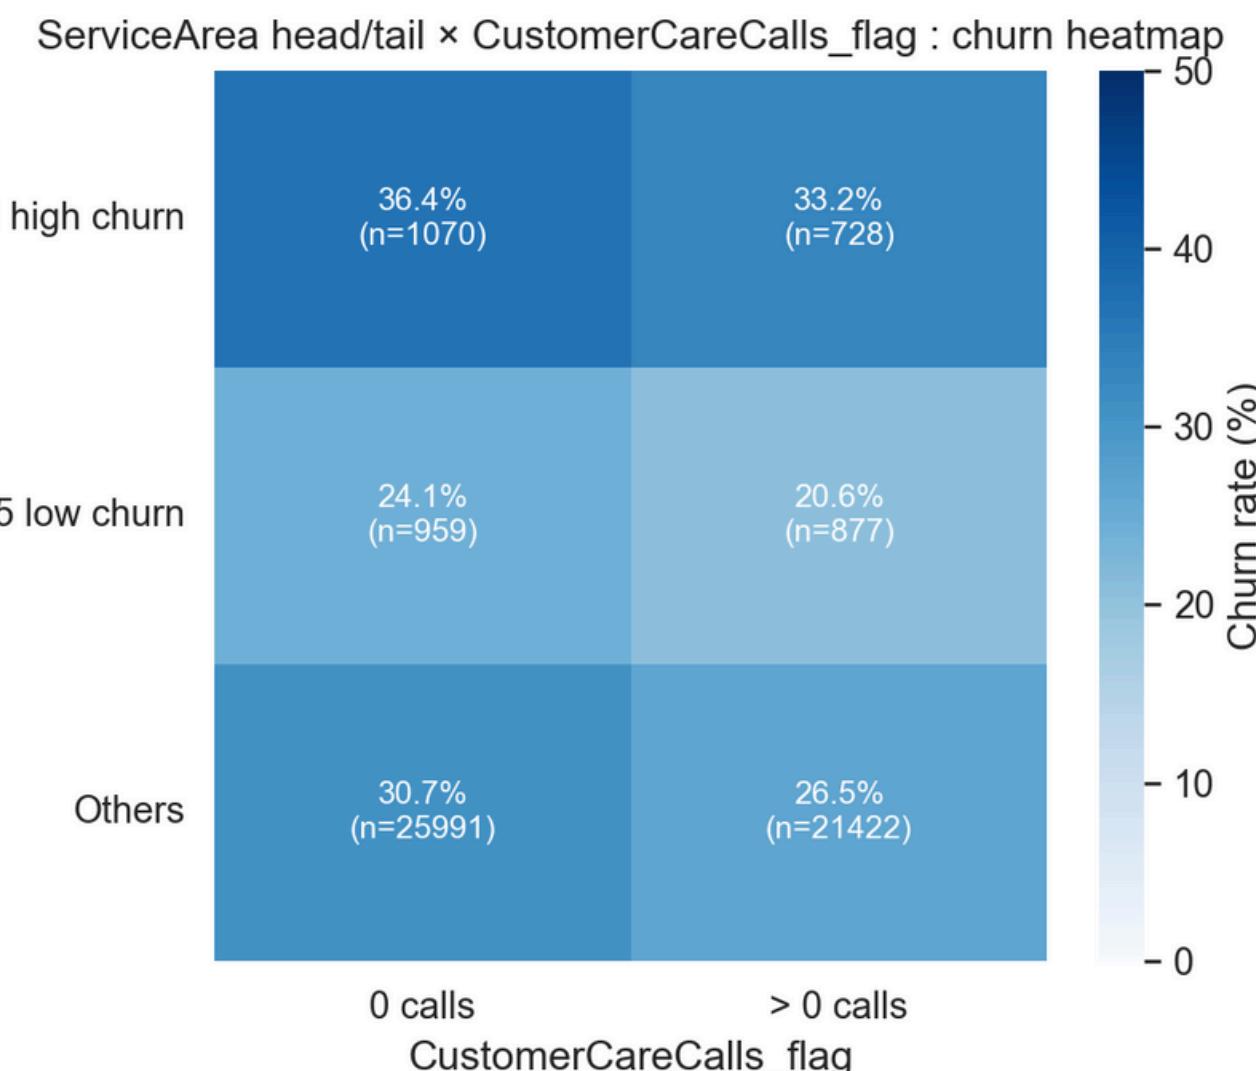
- **Dataset:** telecom churn dataset (tabular, mixed numeric + categorical, imbalanced target)
- **Target:** binary churn label (Churn01)
- **Split strategy:** train/validation/test split done before any threshold learning / encoding fitting
- **Leakage control:** exclude post-treatment / retention-intervention signals (e.g., retention call/accept variables)
- **Evaluation focus:** ranking + decision usefulness under imbalance



Feature Engineering

(Business-friendly + Leakage-safe)

Design principle: make features interpretable and stable; avoid high-cardinality explosion



Examples of engineered features (and why):

- CustomerCareCalls_flag = support/complaint signal (actionable)
- DroppedQuality_flag / DroppedBlockedCalls = network quality proxy
- HighUsage_flag = usage pressure (threshold learned on train)
- HandsetPrice_q = price binning + missing as category (robust)
- HighRiskServiceArea_flag = compress high-cardinality area into risk flag

Models Compared

Four Families

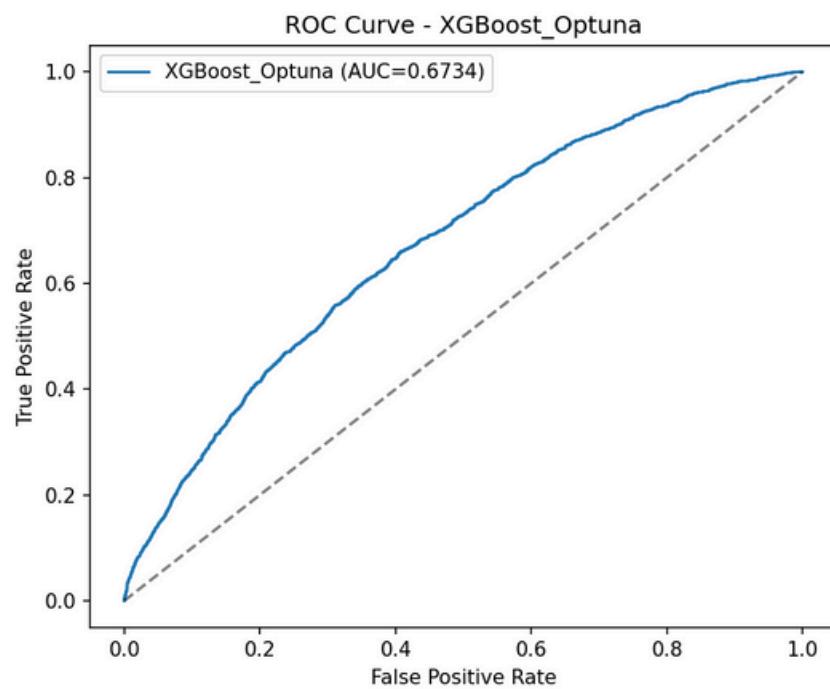
- **Baseline:** Logistic Regression (interpretable, strong calibration baseline)
- **Tree/GBM:** Decision Tree, RandomForest, LightGBM, XGBoost (nonlinear + interactions)
- **Deep Learning:** Dense MLP, Embedding NN, Wide & Deep; imbalance strategies (strong weighting, focal loss)
- **Ensemble:** OOF Stacking over diverse base learners



Experimental Design & Metrics

How I Compare Fairly

- Same train/val/test split across all models
- Hyperparameter tuning limited to validation (no test peeking)
- **Primary metrics:** ROC-AUC and PR-AUC (imbalance-aware)
- **Decision metrics:** F1 / Recall at tuned threshold; and threshold is chosen on validation



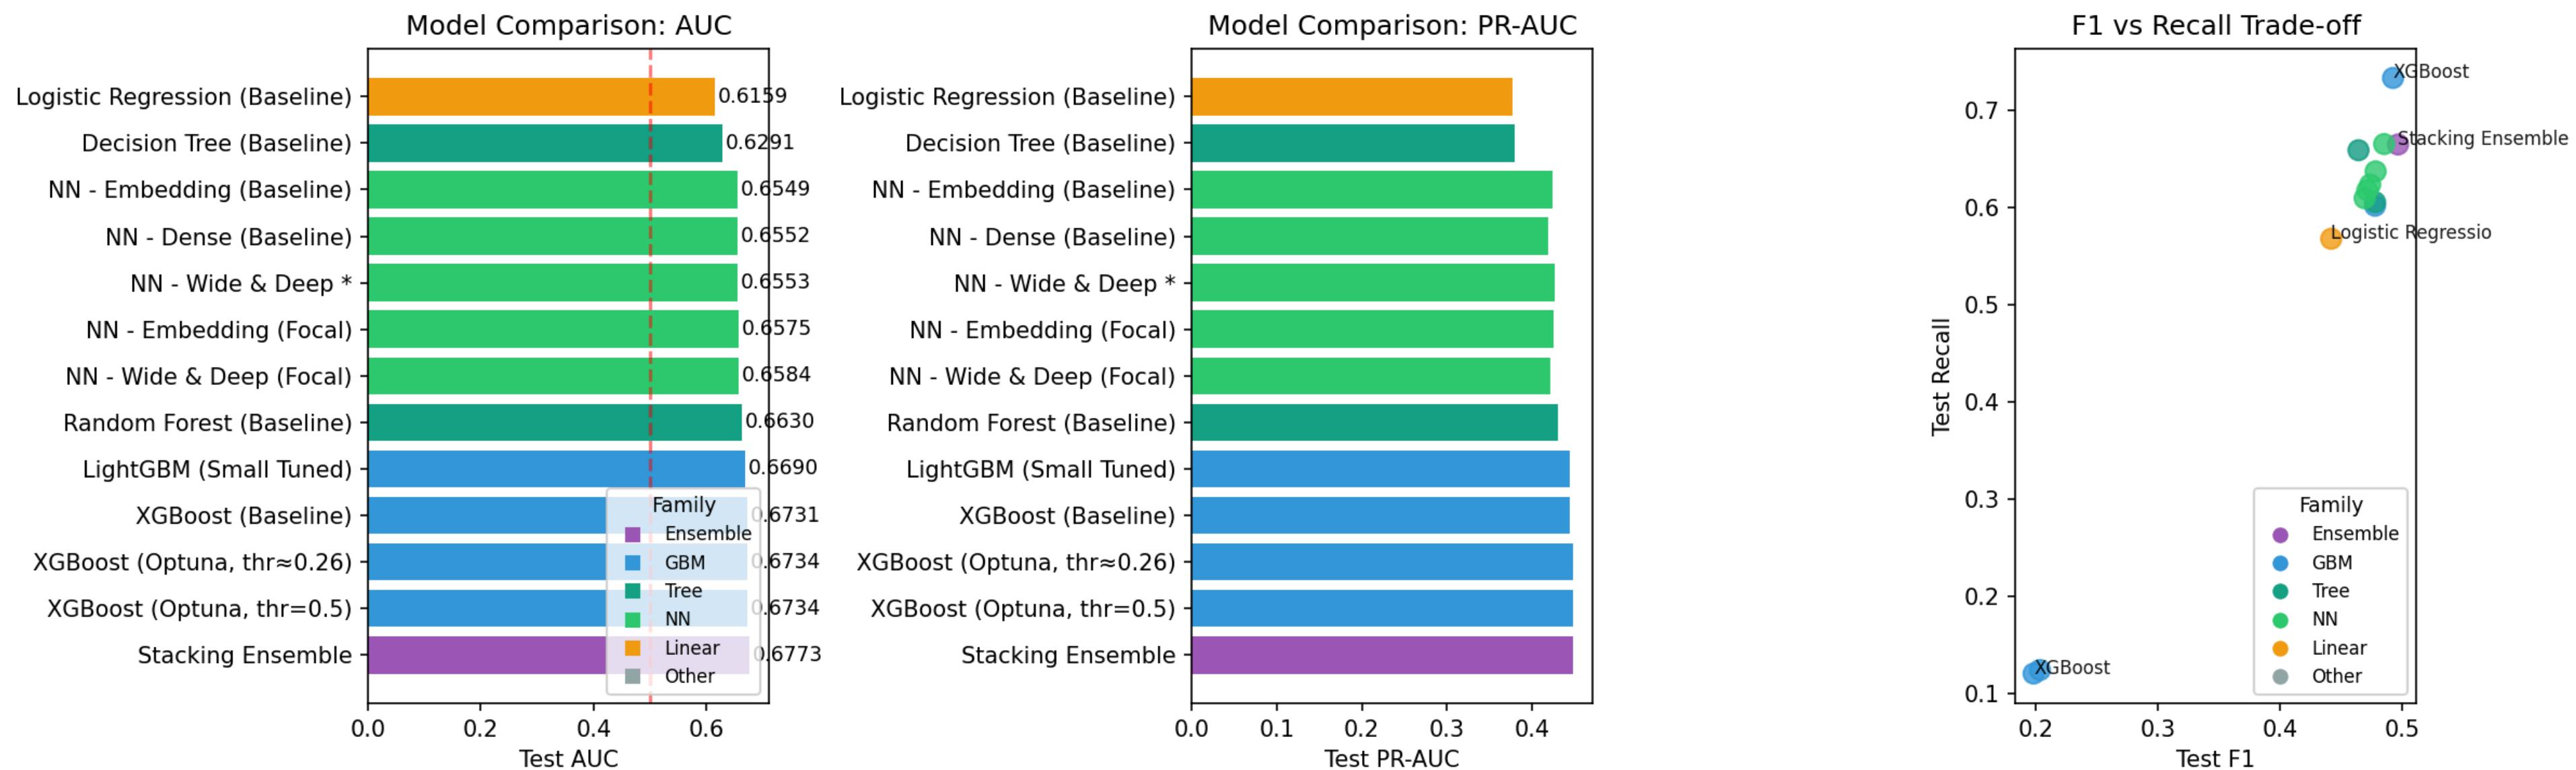
Metric	Why it matters
ROC-AUC	Ranking quality: how well the model orders churners above non-churners across all thresholds.
PR-AUC	Positive-class focus: more informative under class imbalance; emphasizes precision/recall on churners.
Confusion matrix	Decision trade-off: shows FP/FN at a chosen threshold → links directly to business costs and intervention policy.



Main Results

Leaderboard + Key Takeaway

Model Leaderboard

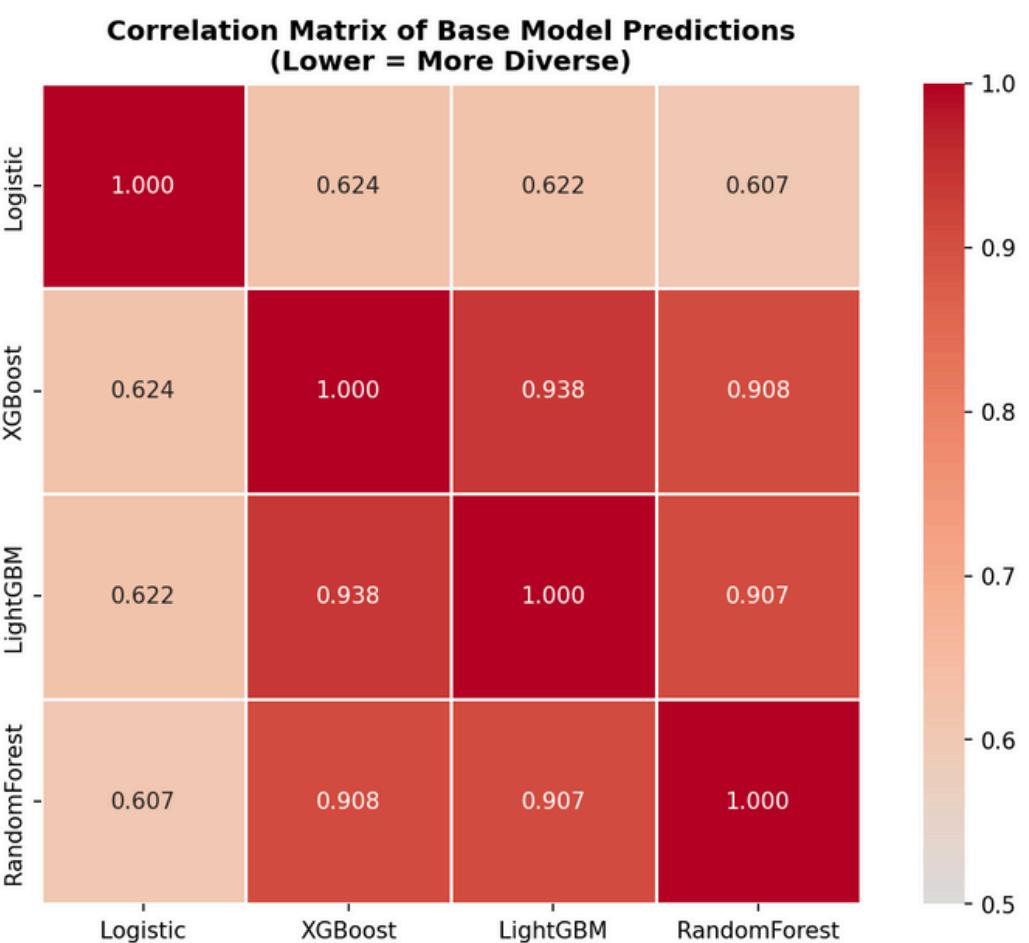


Advanced Modeling Beyond GBM

Tabular Deep Learning + Leakage-safe Stacking

- Controlled study: 7 NN variants under the same data split + training protocol(dense / embedding / wide&deep / deeper / focal / weighting)
- Best NN: Wide & Deep – Val AUC 0.6356, Test AUC 0.6553, PR-AUC 0.4259
- Imbalance handling: used class-weighted BCE (`pos_weight`); compared with stronger weighting and focal loss ablations
- Takeaway: Embeddings + wide component help, but gain is limited vs top GBMs on this tabular dataset

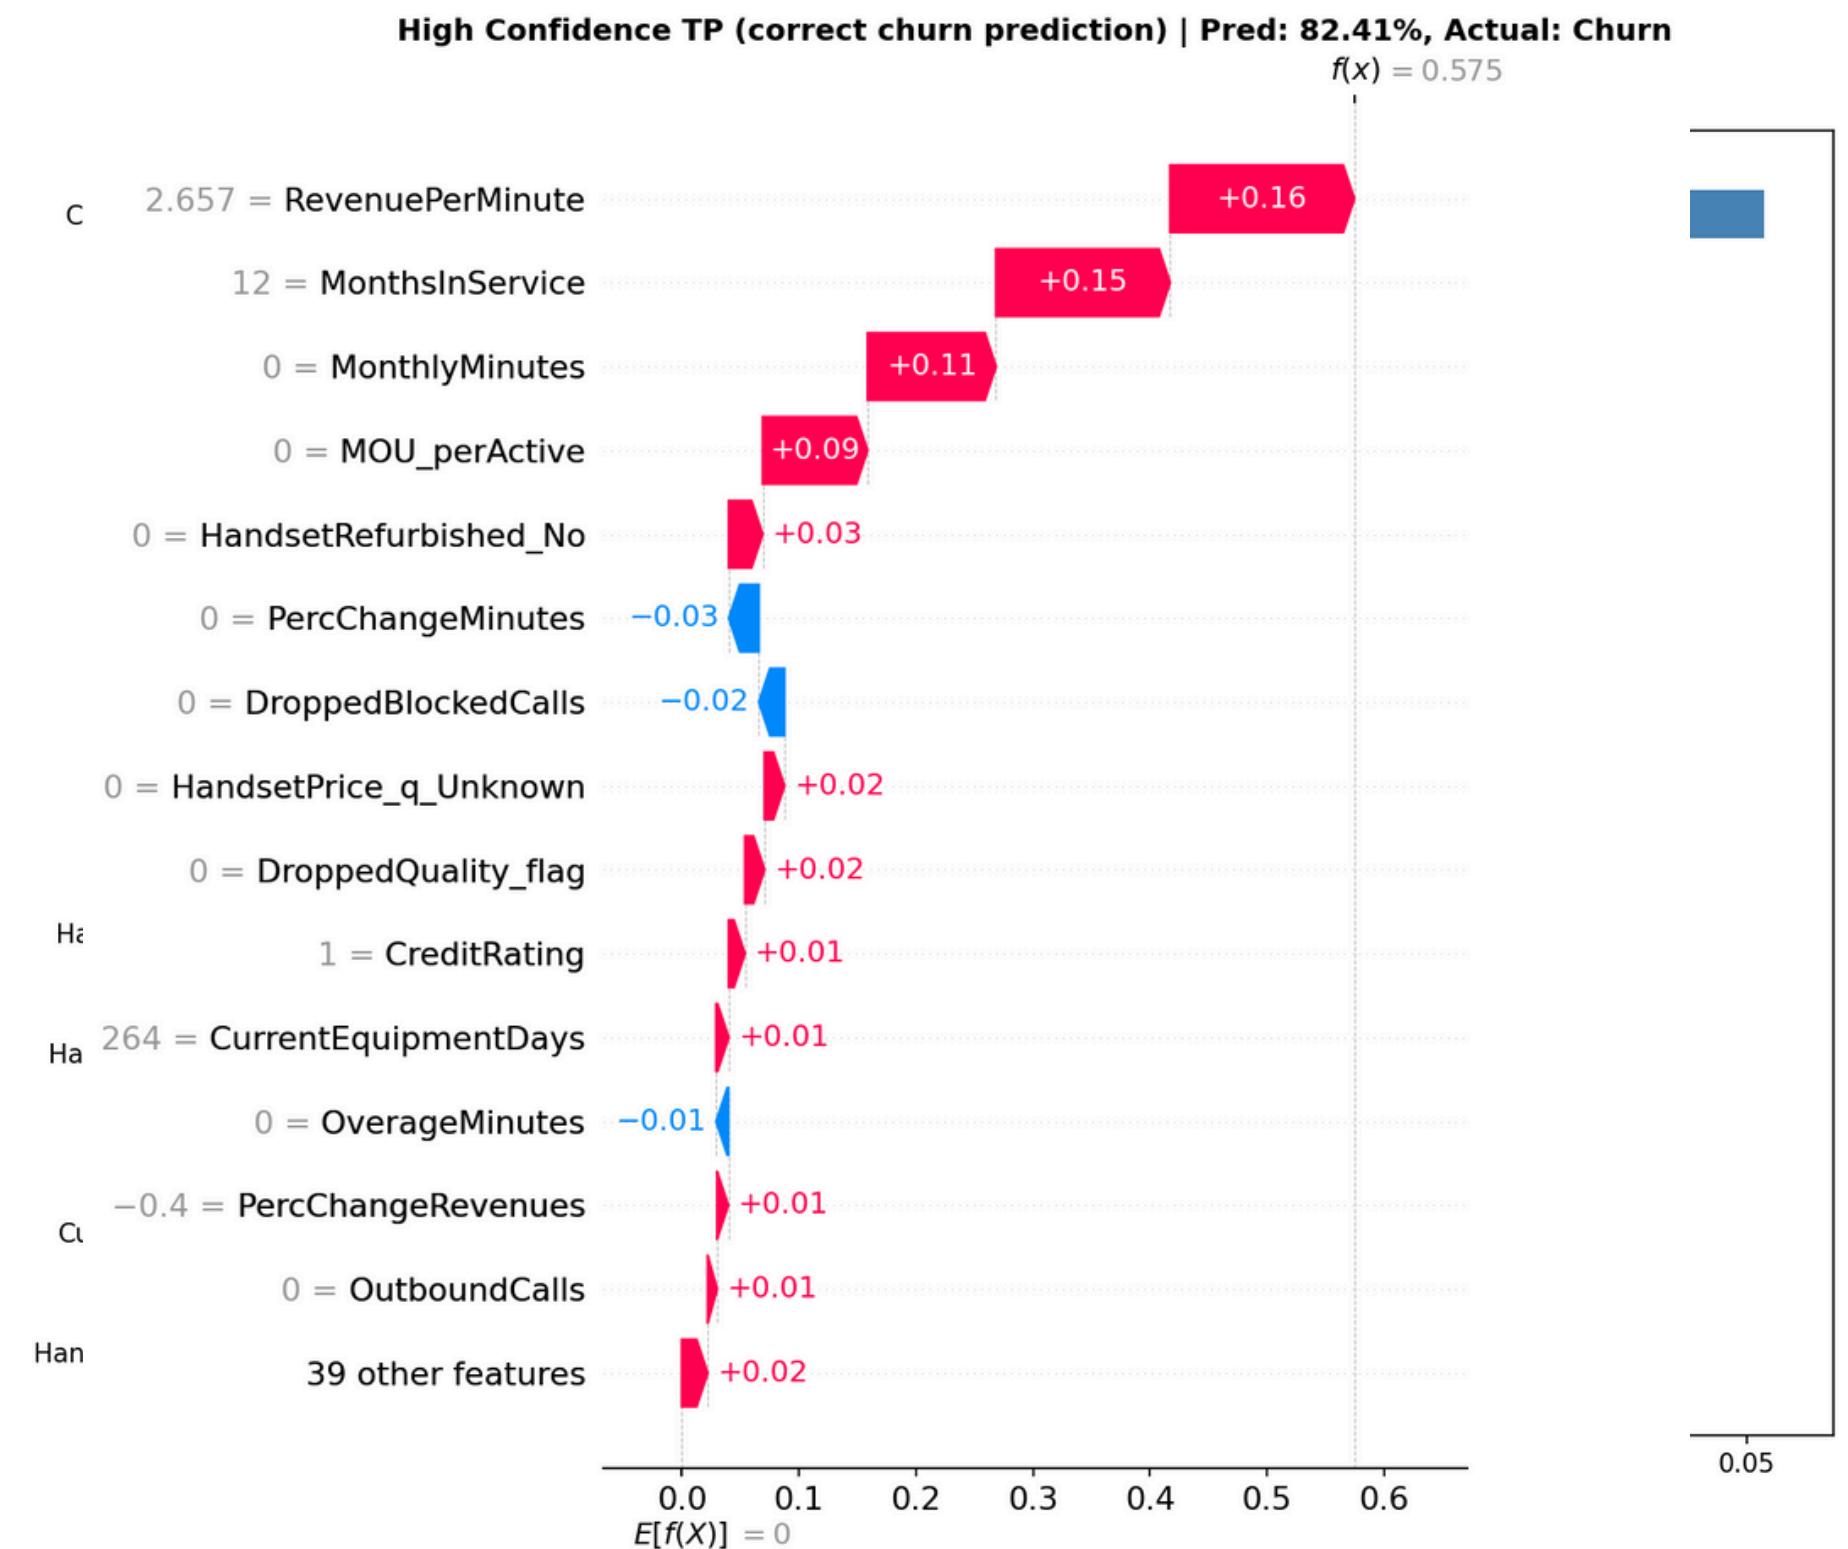
- **Why stacking:** combine complementary errors across model families
- Evidence of diversity: avg prediction correlation ≈ 0.768 (not identical signals)
- Negative result: naïve averaging did not beat the best single model (important sanity check)
- Leakage-safe method: meta-learner trained on out-of-fold (OOF) base predictions
- Result: Stacking Test AUC 0.6773 vs best single XGB 0.6734 (+0.58% relative lift)



Explainability: SHAP Insights

From “Black Box” to Action

- Global SHAP shows top churn drivers:
**CurrentEquipmentDays (largest),
PercChangeMinutes, MonthsInService,
DroppedBlockedCalls, RevenuePerMinute,
MonthlyMinutes, etc**
- Use SHAP dependence plots to show
nonlinear risk patterns (e.g., device age /
usage change effects)
- SHAP is **associational**, not causal →
interventions should be validated with
experiments



Business Case Study

Why calibrated probabilities matter

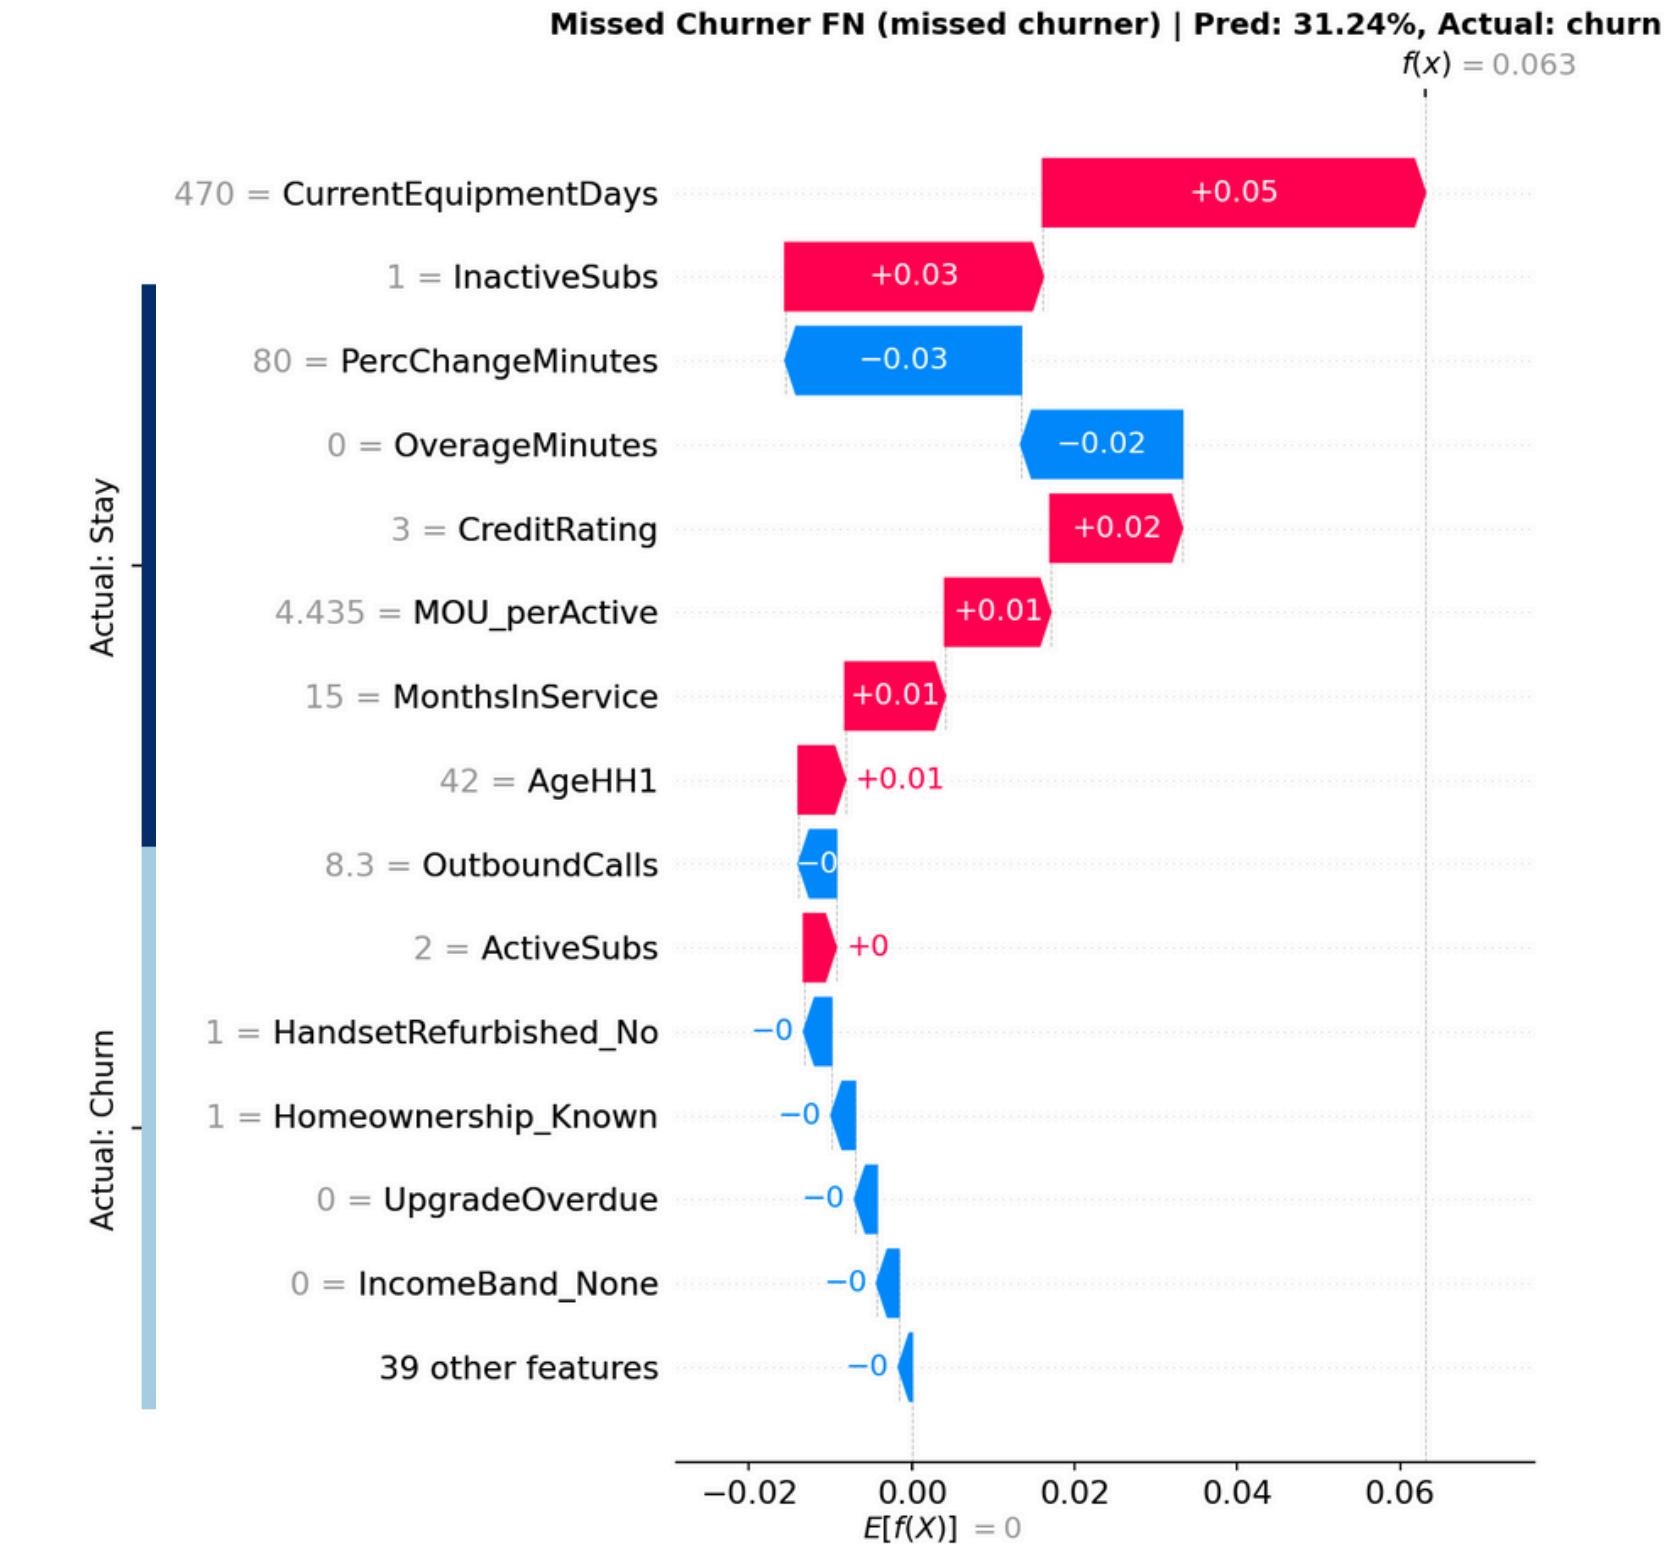
- Decision rule (expected value):
 - **Expected Loss = $P(\text{churn}) \times \text{LTV}$**
 - Retain only if: $P(\text{churn}) \times \text{LTV} > \text{Cost}(\text{retention})$
- Example assumption:
 - Cost = \$200, LTV = \$1000 → break-even probability = 0.20
- Why calibration matters:
 - If model says 0.80 but true churn rate is 0.10, you over-spend and get negative ROI (you're “buying insurance” that's too expensive).



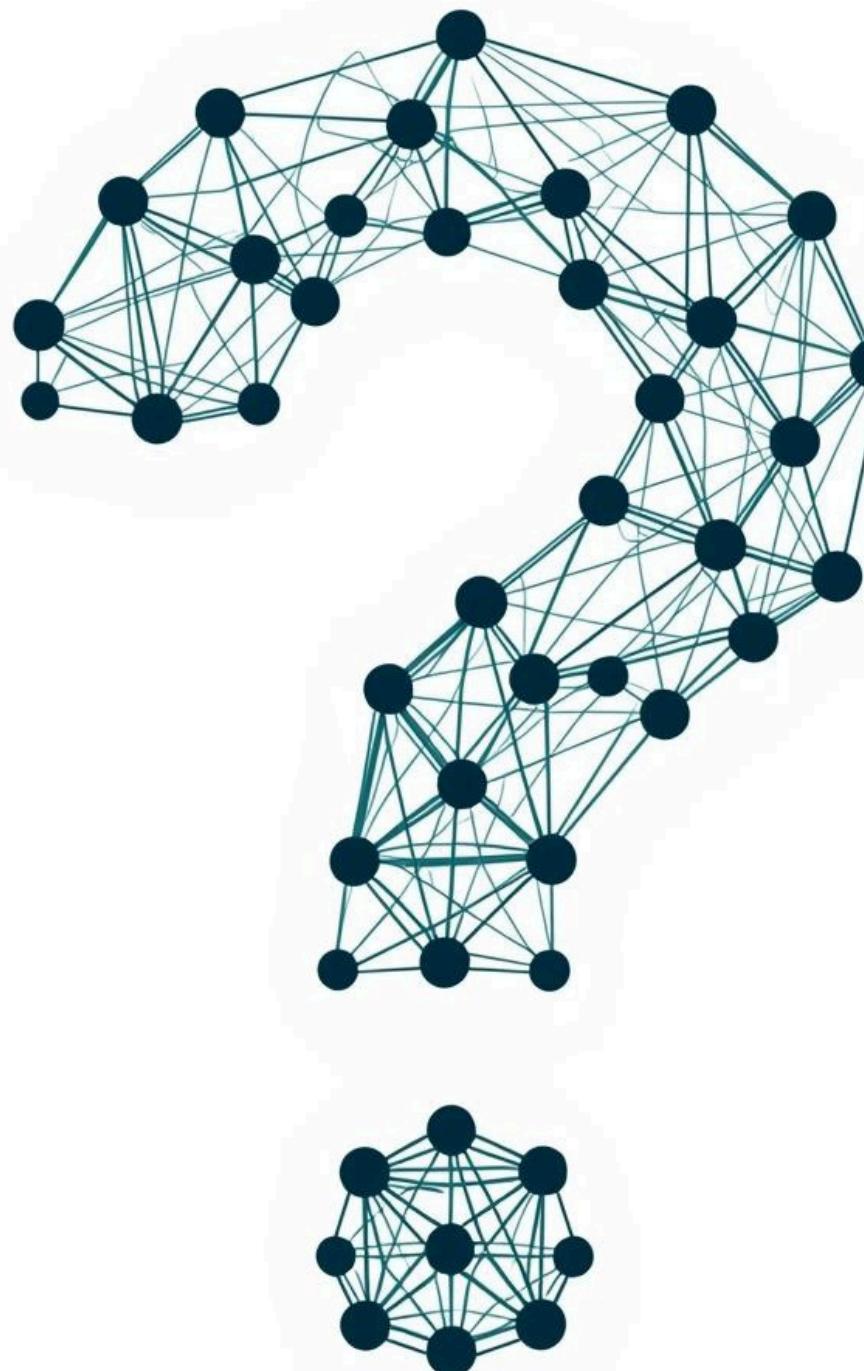
Error Analysis & Decisioning

Threshold + Confusion Matrix + Next Steps

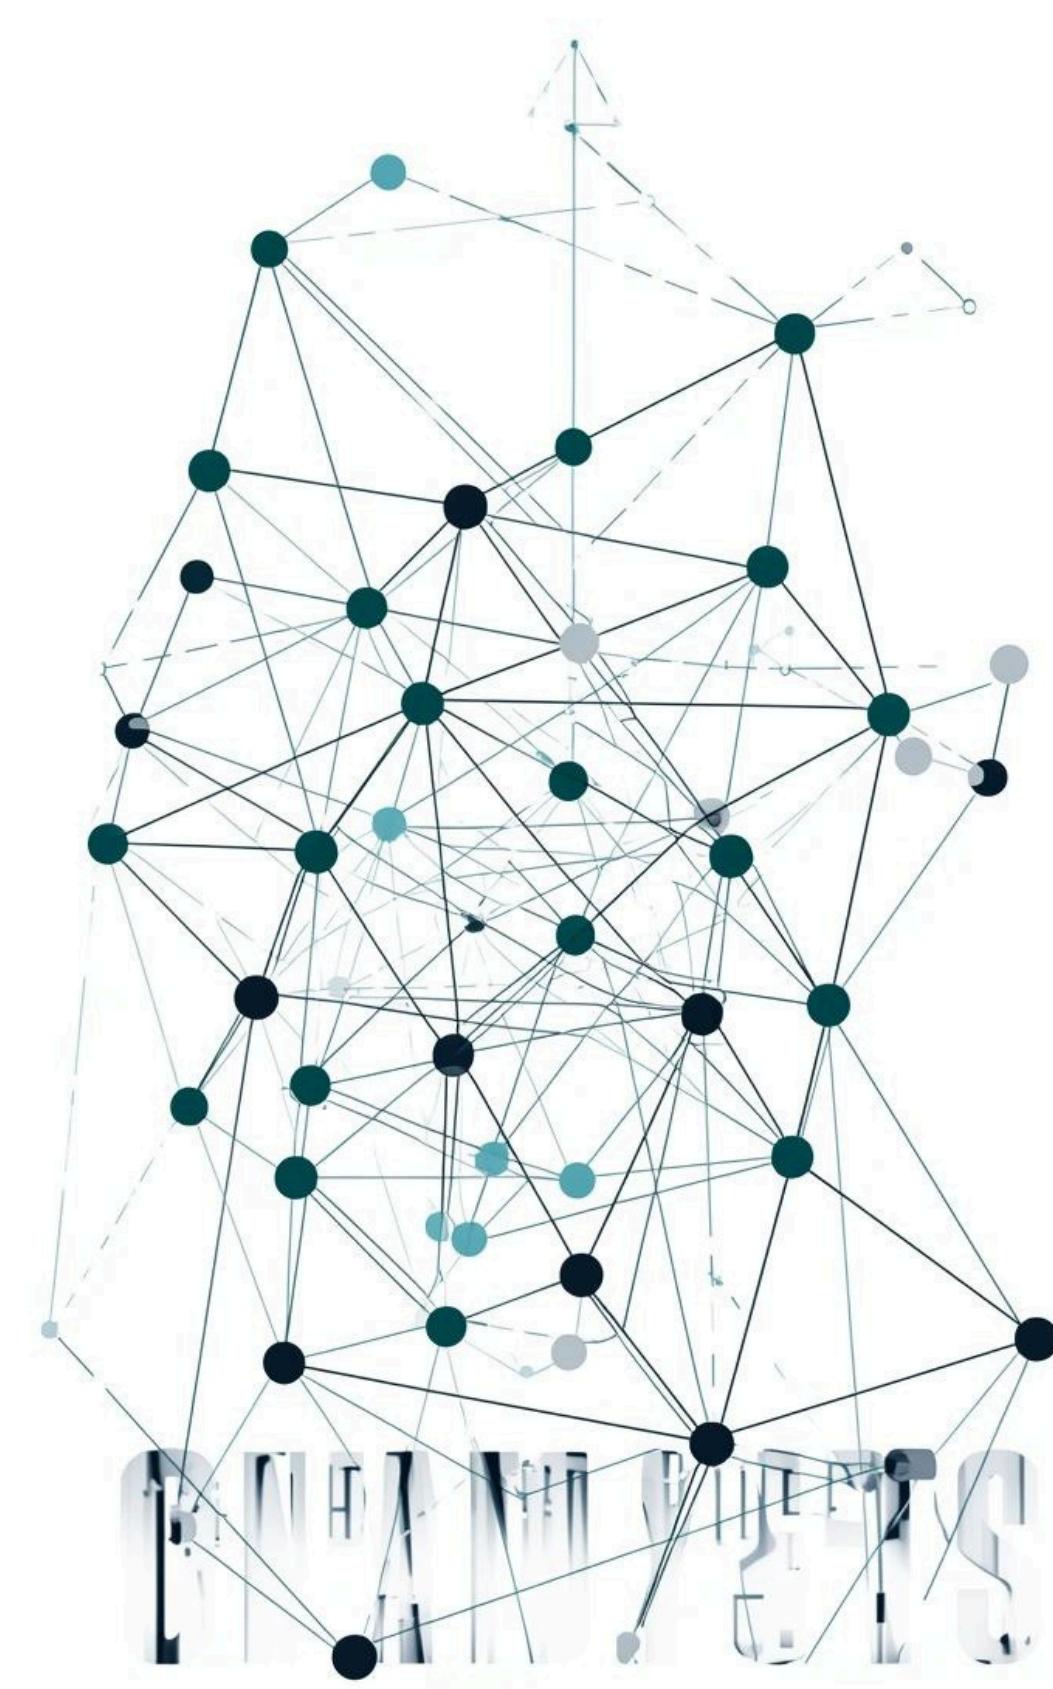
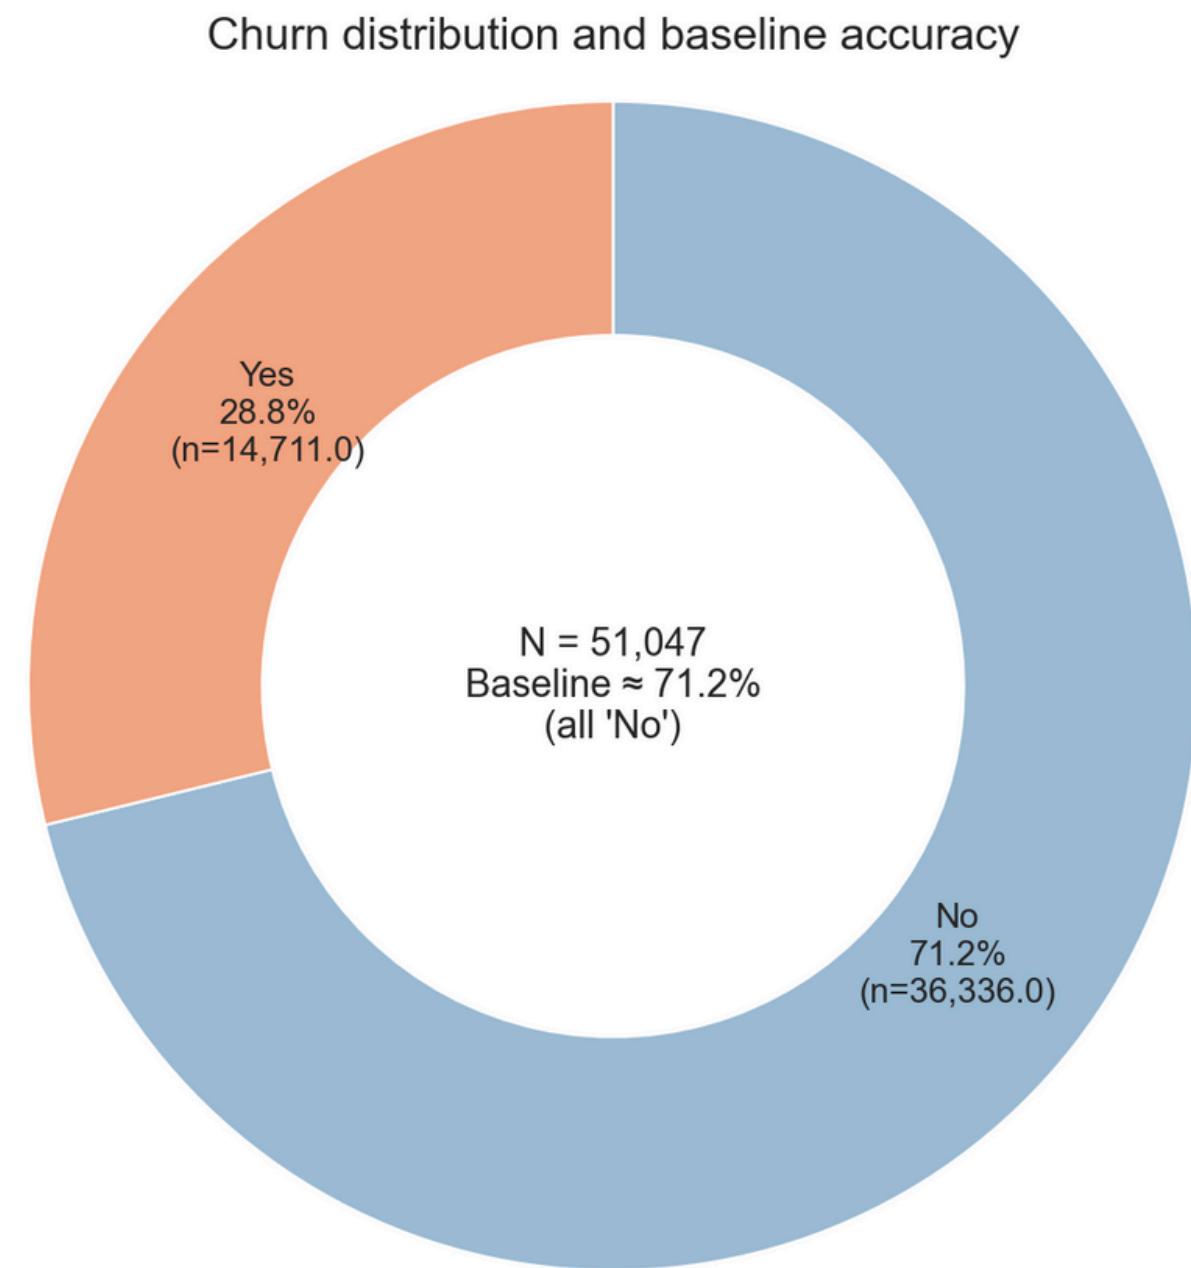
- Default threshold 0.5 can miss churners (example confusion matrix shows churn recall is low at 0.5)
- In my paper-style evaluation, I also tune threshold to hit business goals (e.g., threshold ≈ 0.2577 giving churn recall ≈ 0.734)
- Show one FN SHAP waterfall: competing signals + threshold too high \rightarrow action: adjust threshold, add features, or segment strategy
- Next steps: calibration check, uplift testing, monitoring drift by key driver



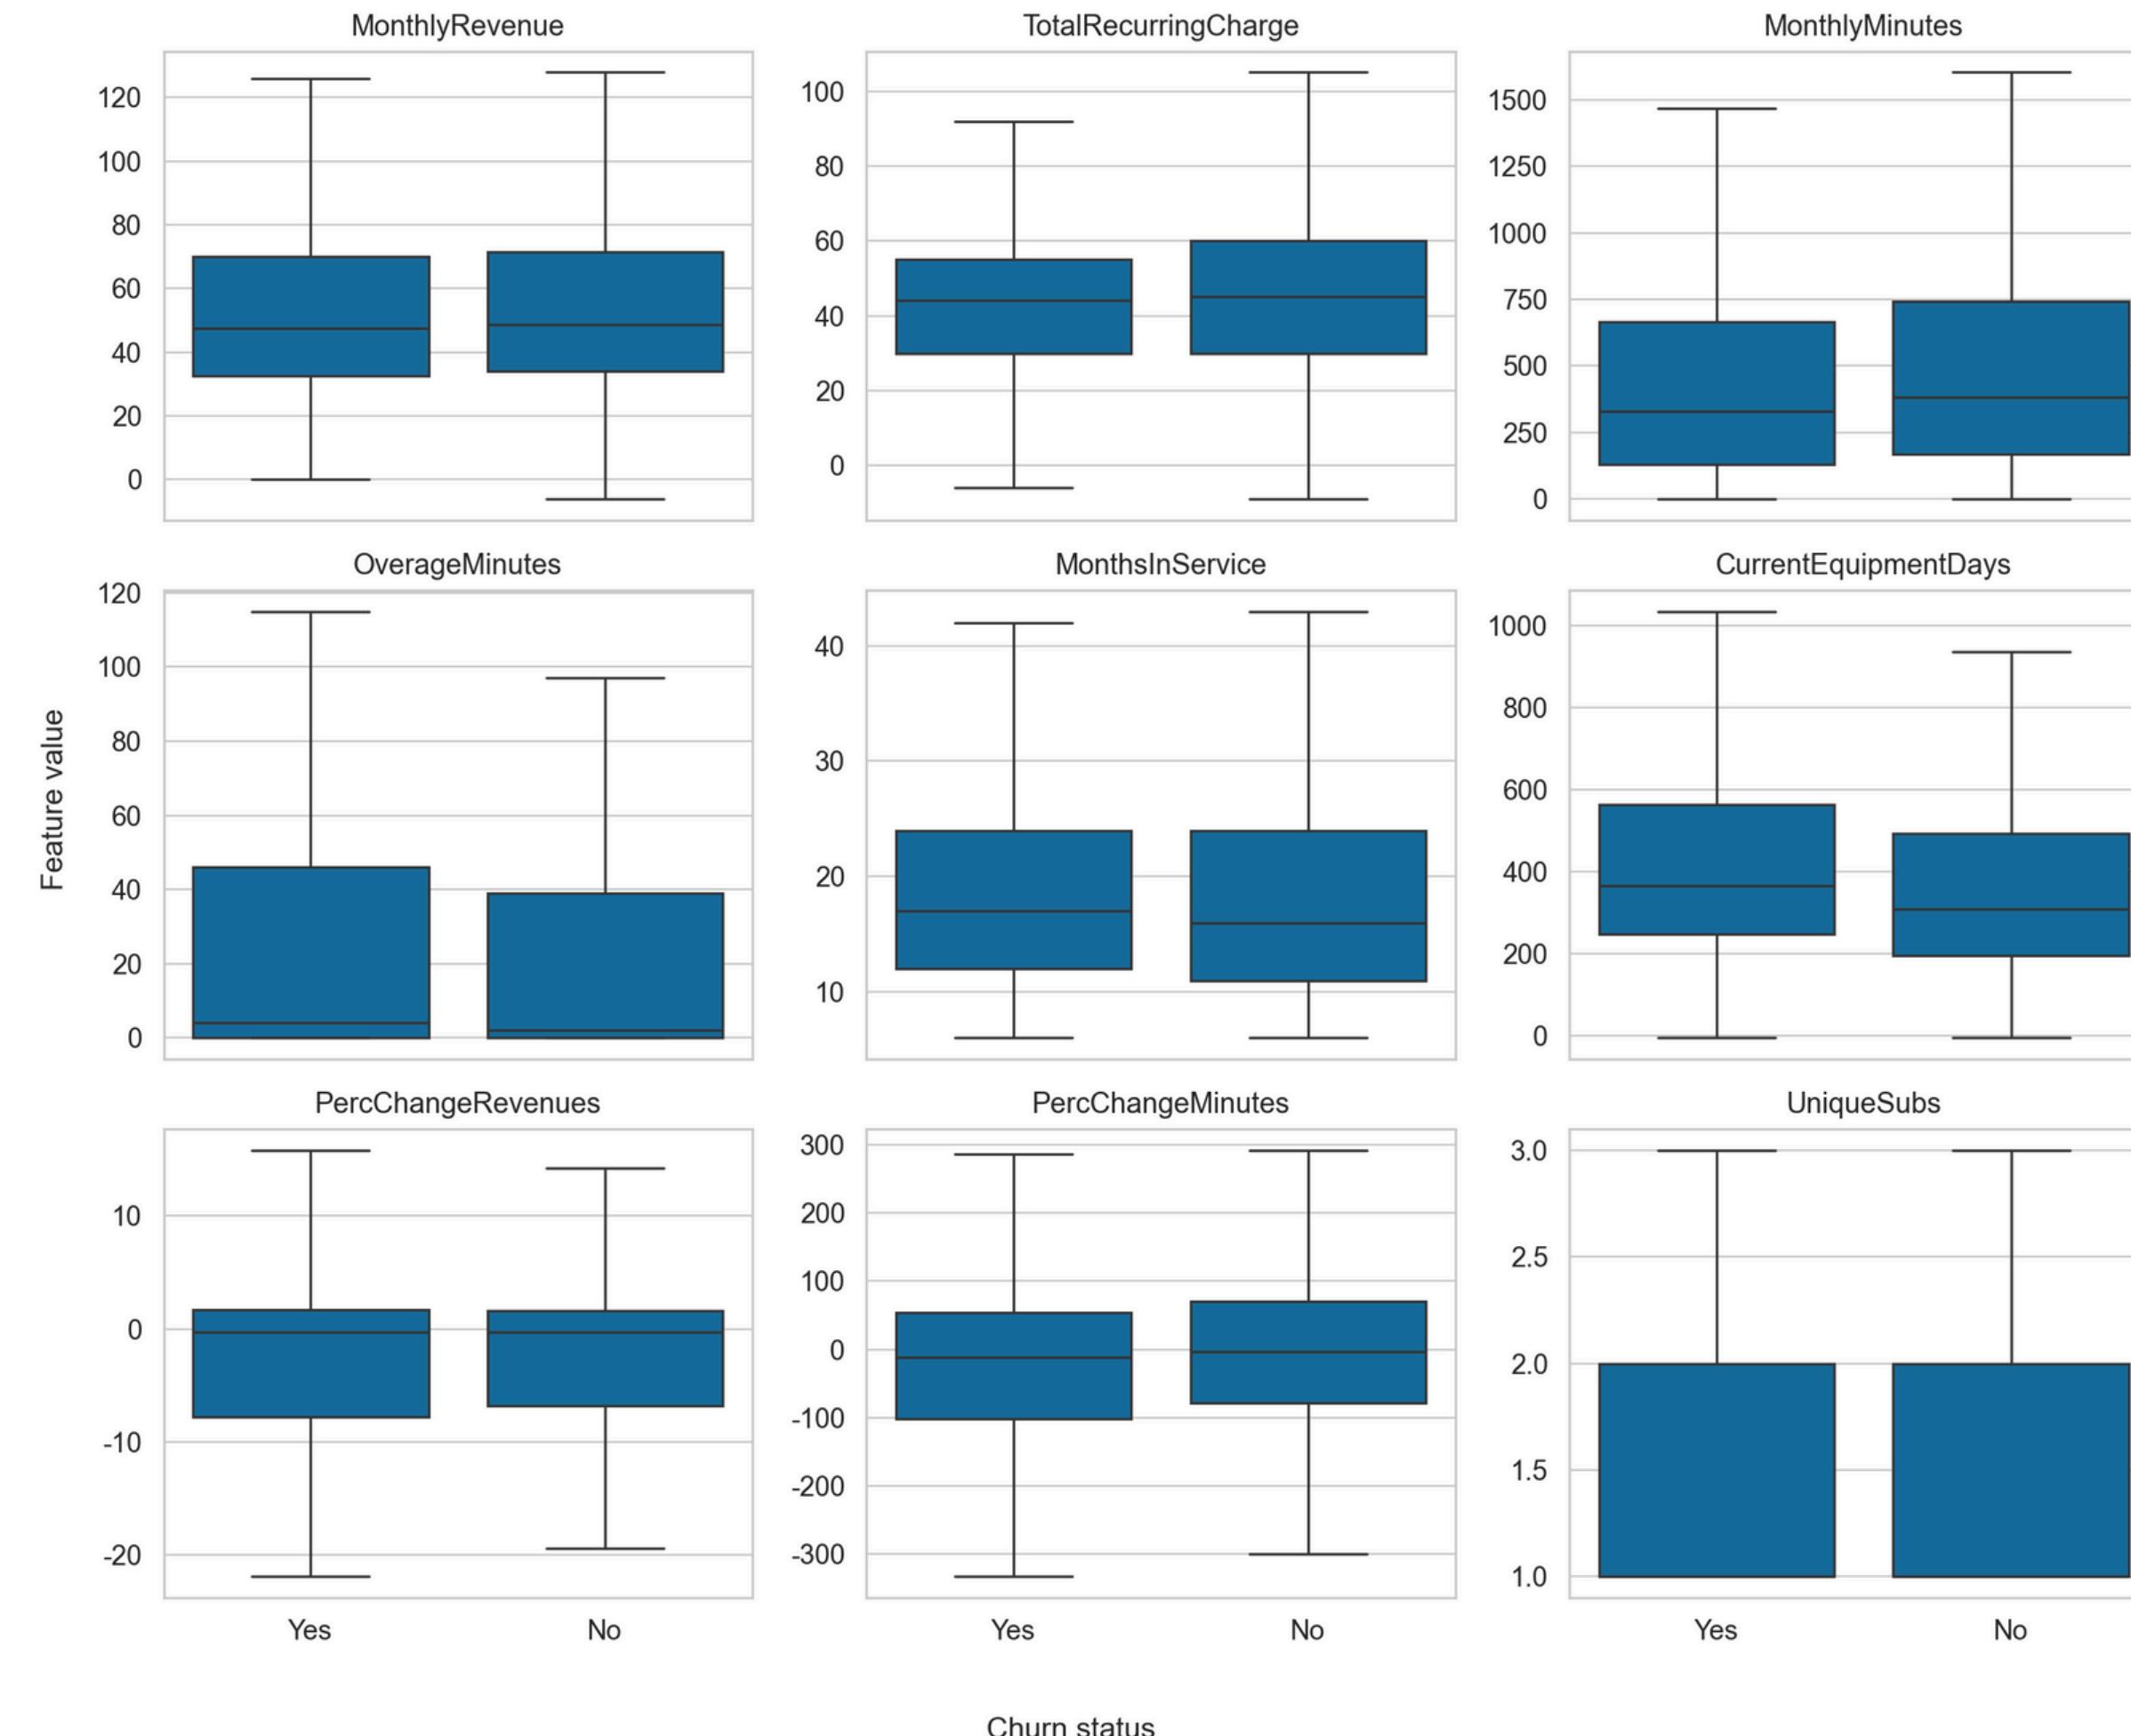
Q&A



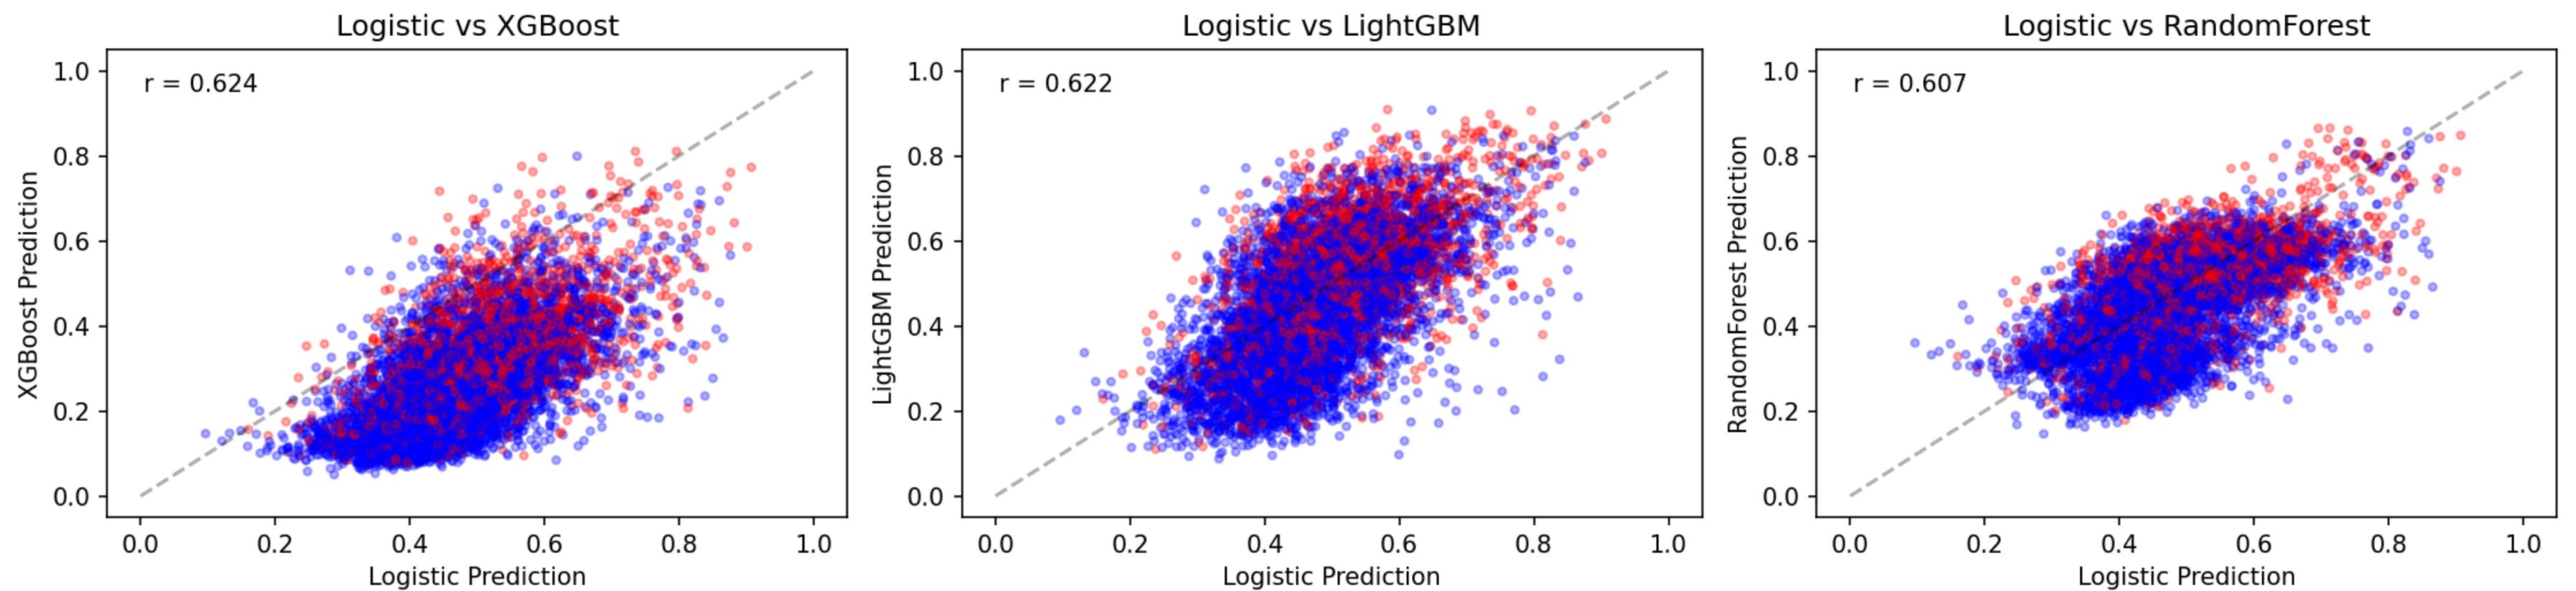
Backup Slide Menu

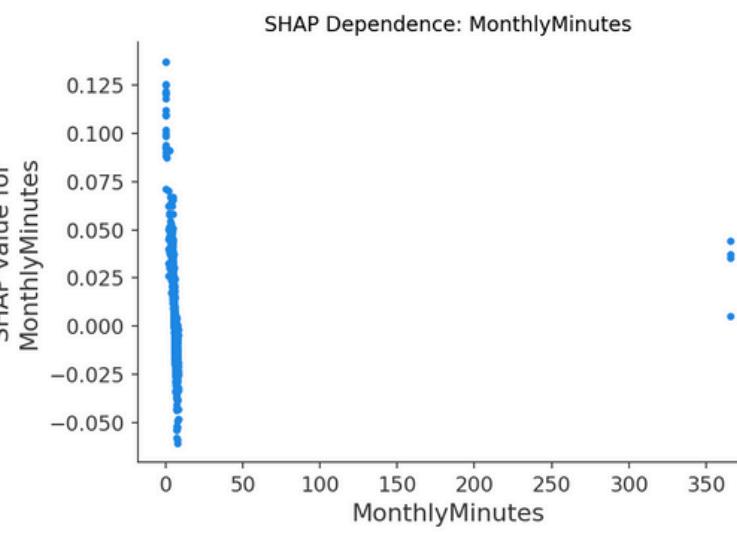
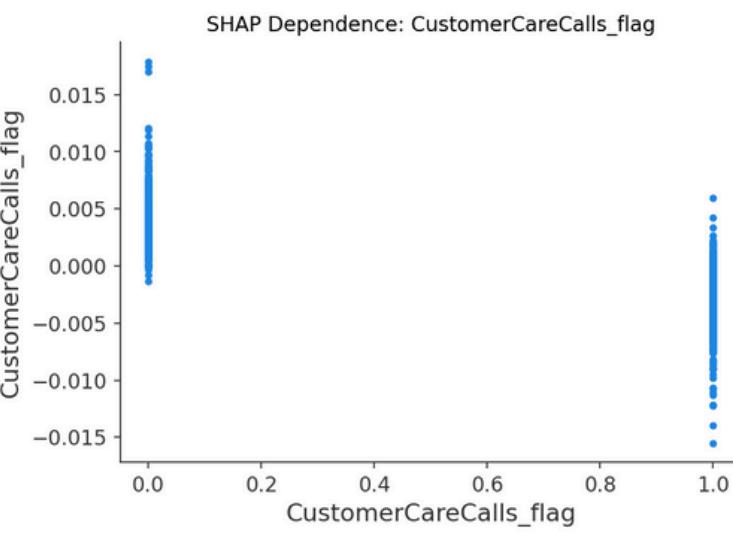
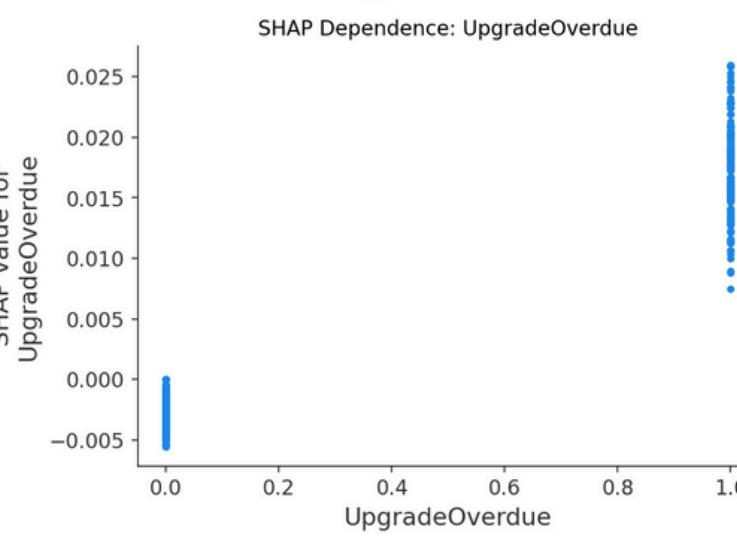
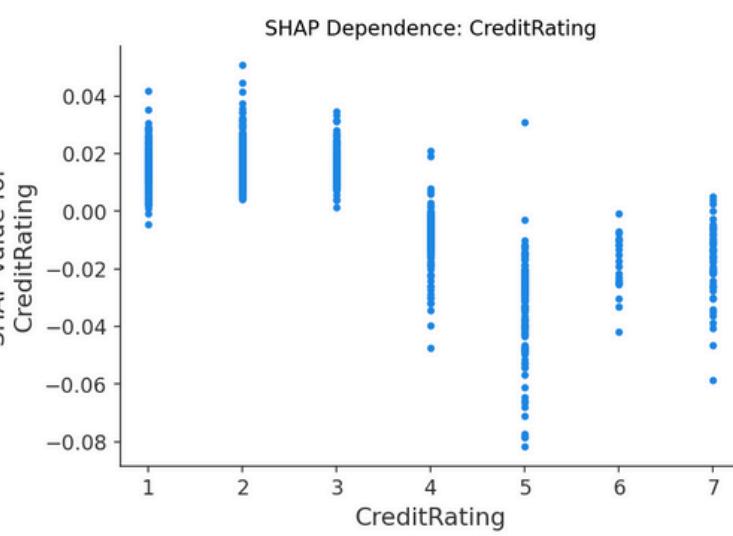
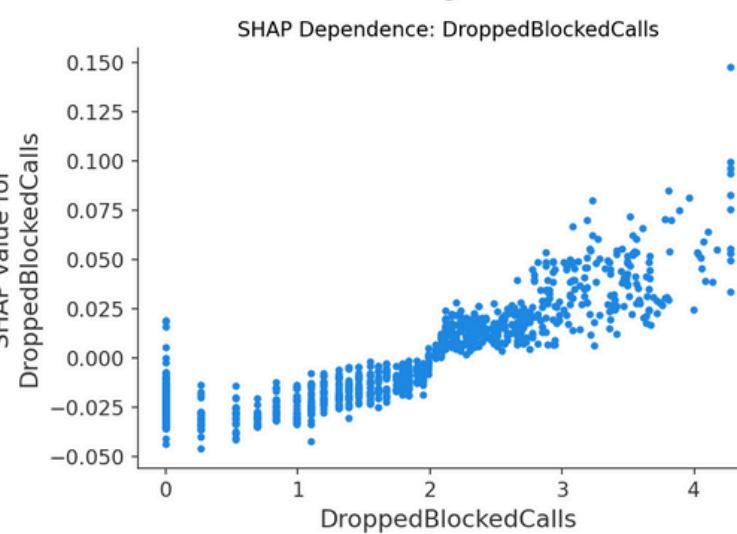
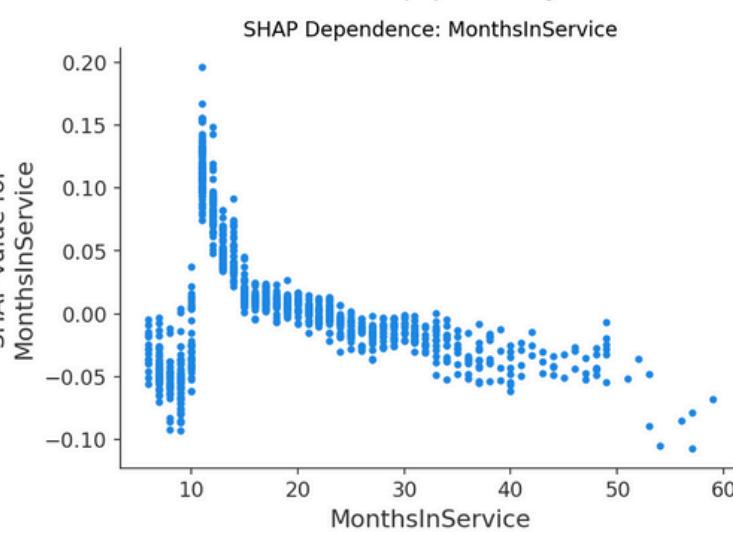
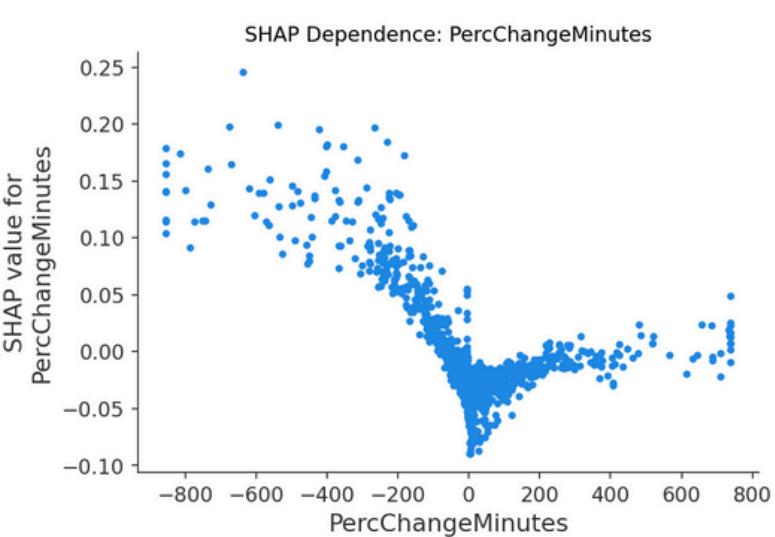
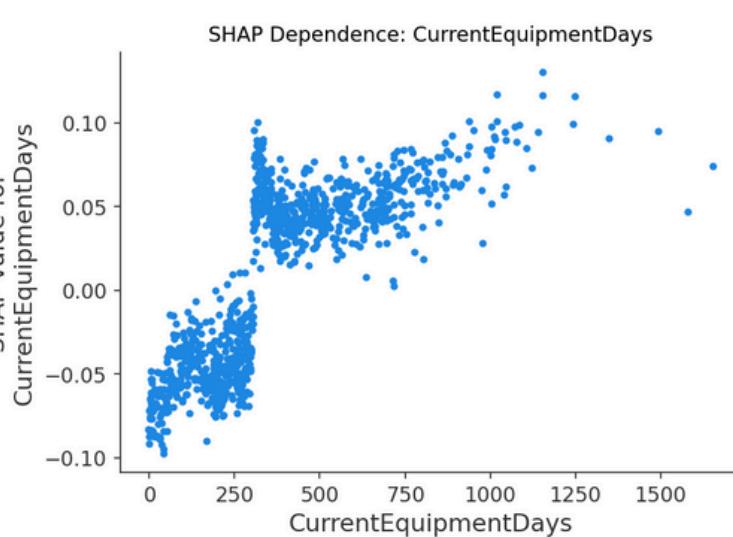


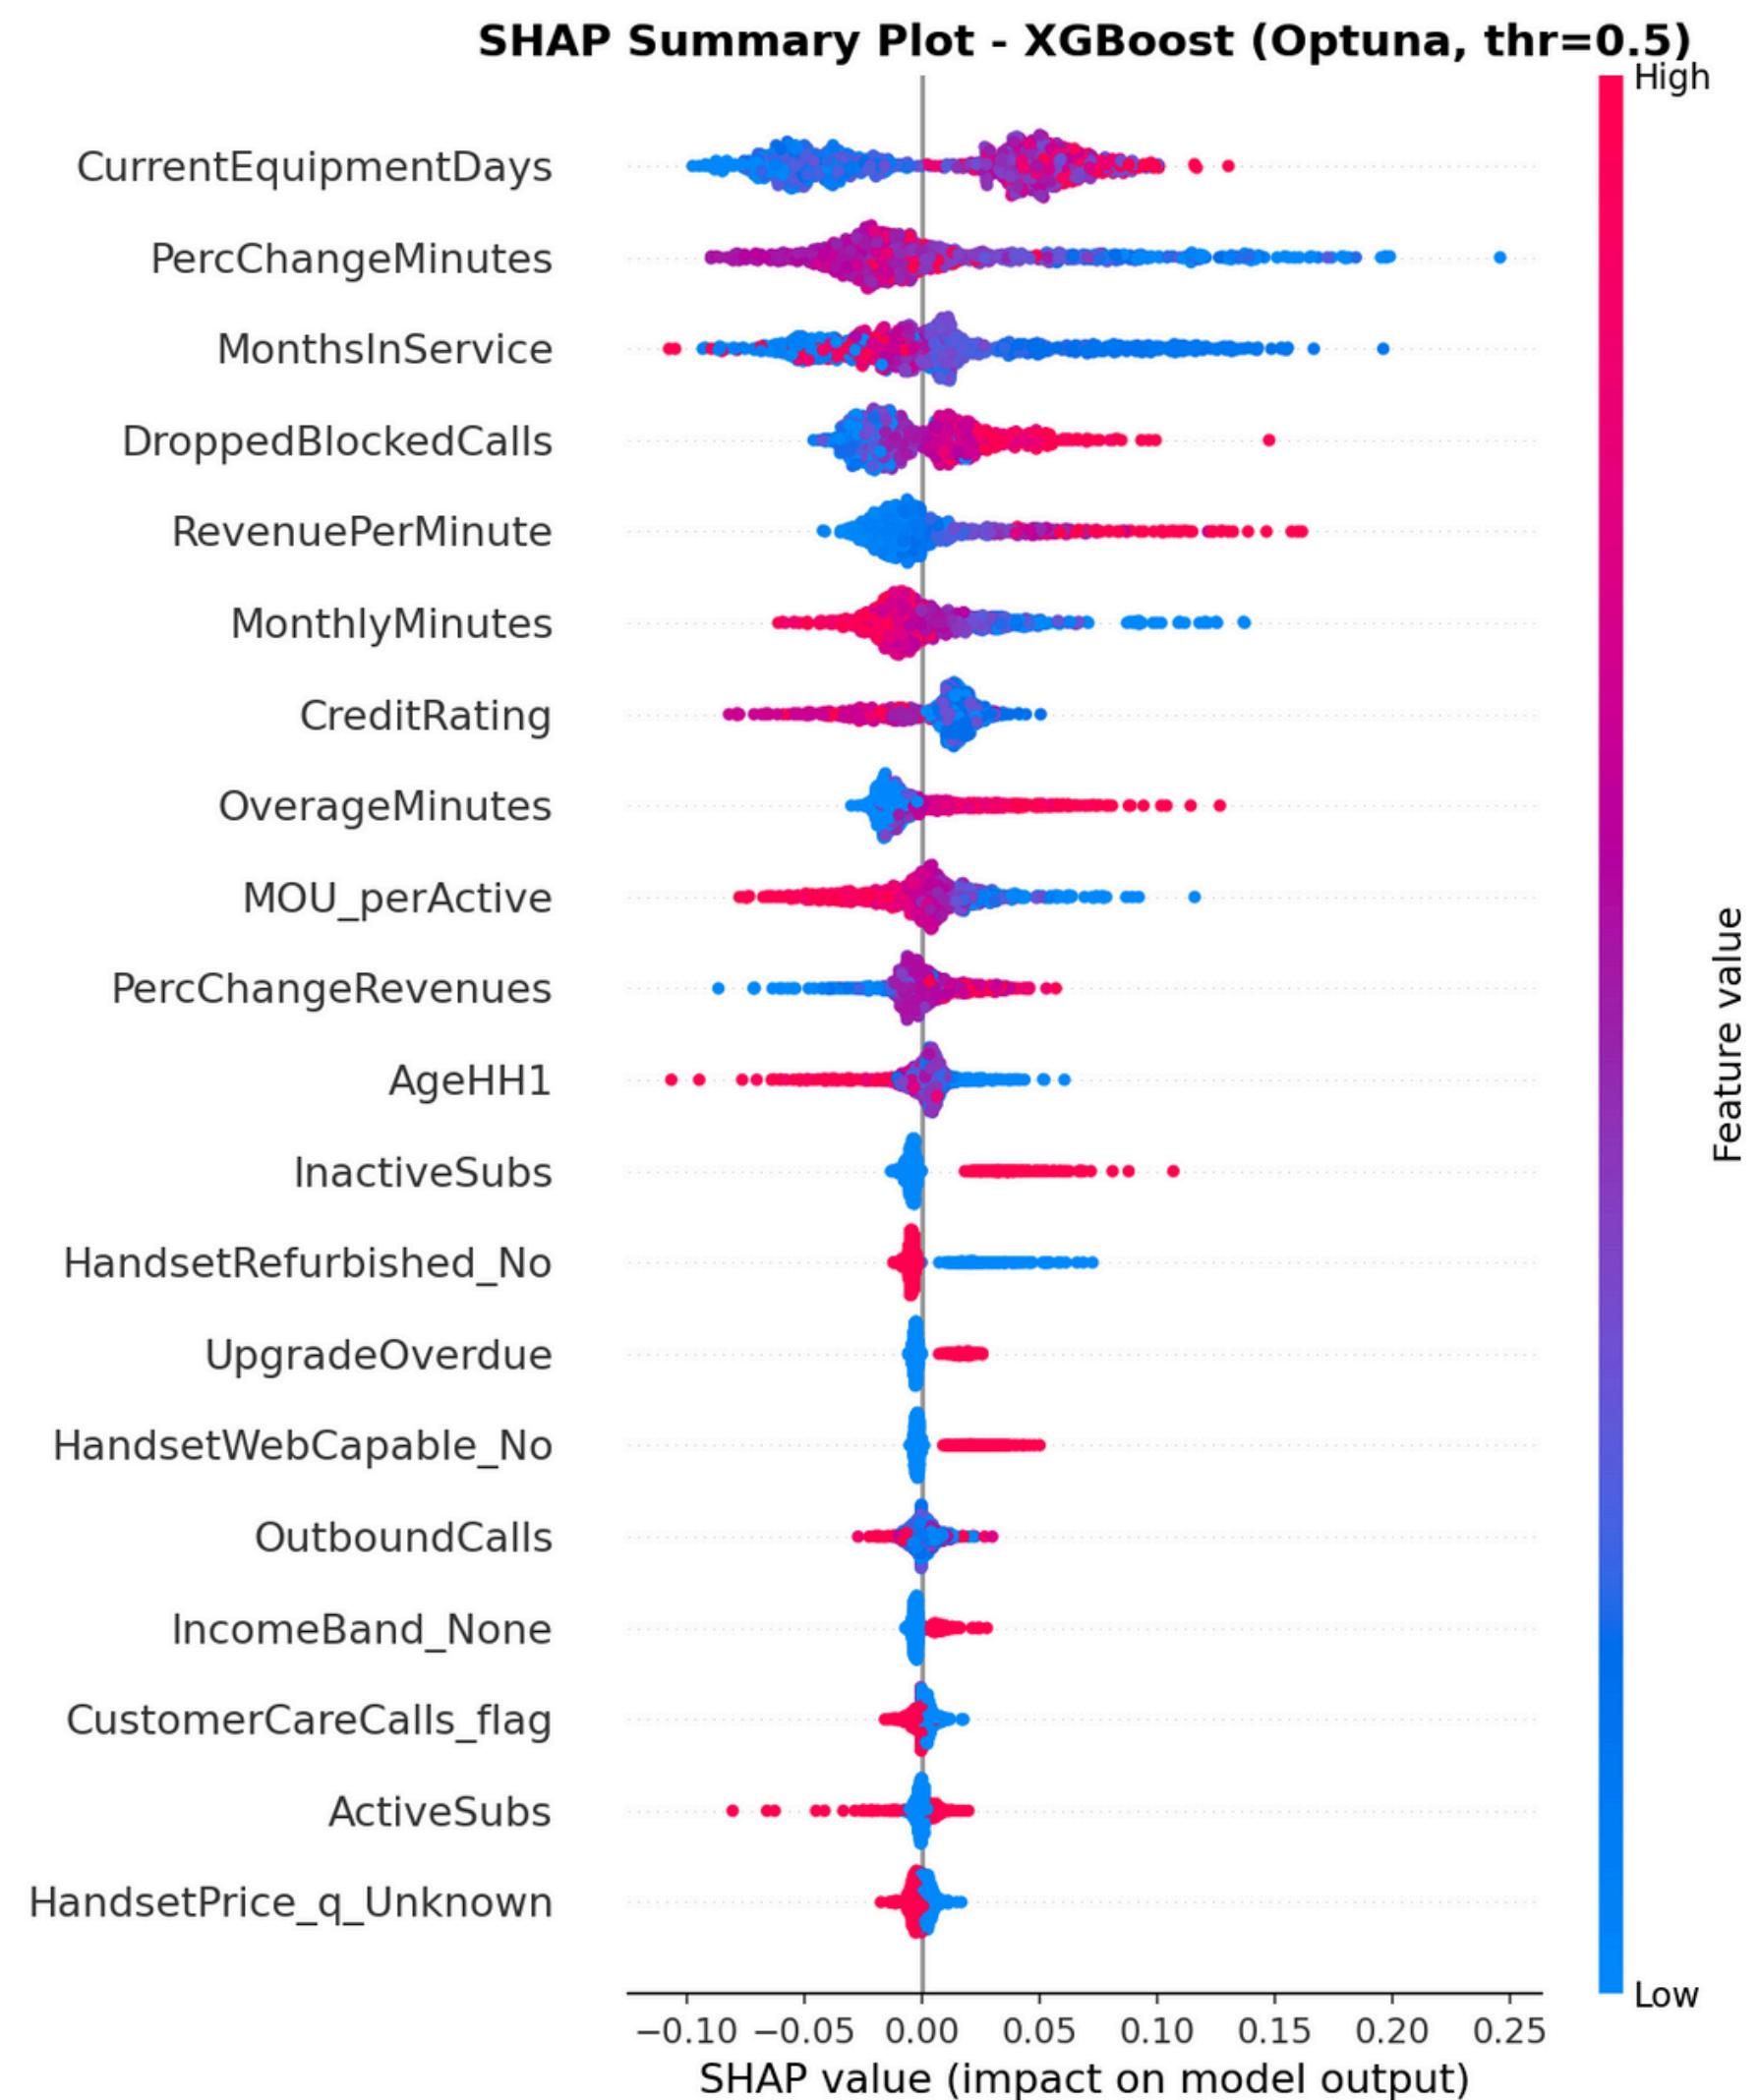
Distribution of key numeric features by Churn



Base Model Predictions (Red=Churn, Blue=Stay)





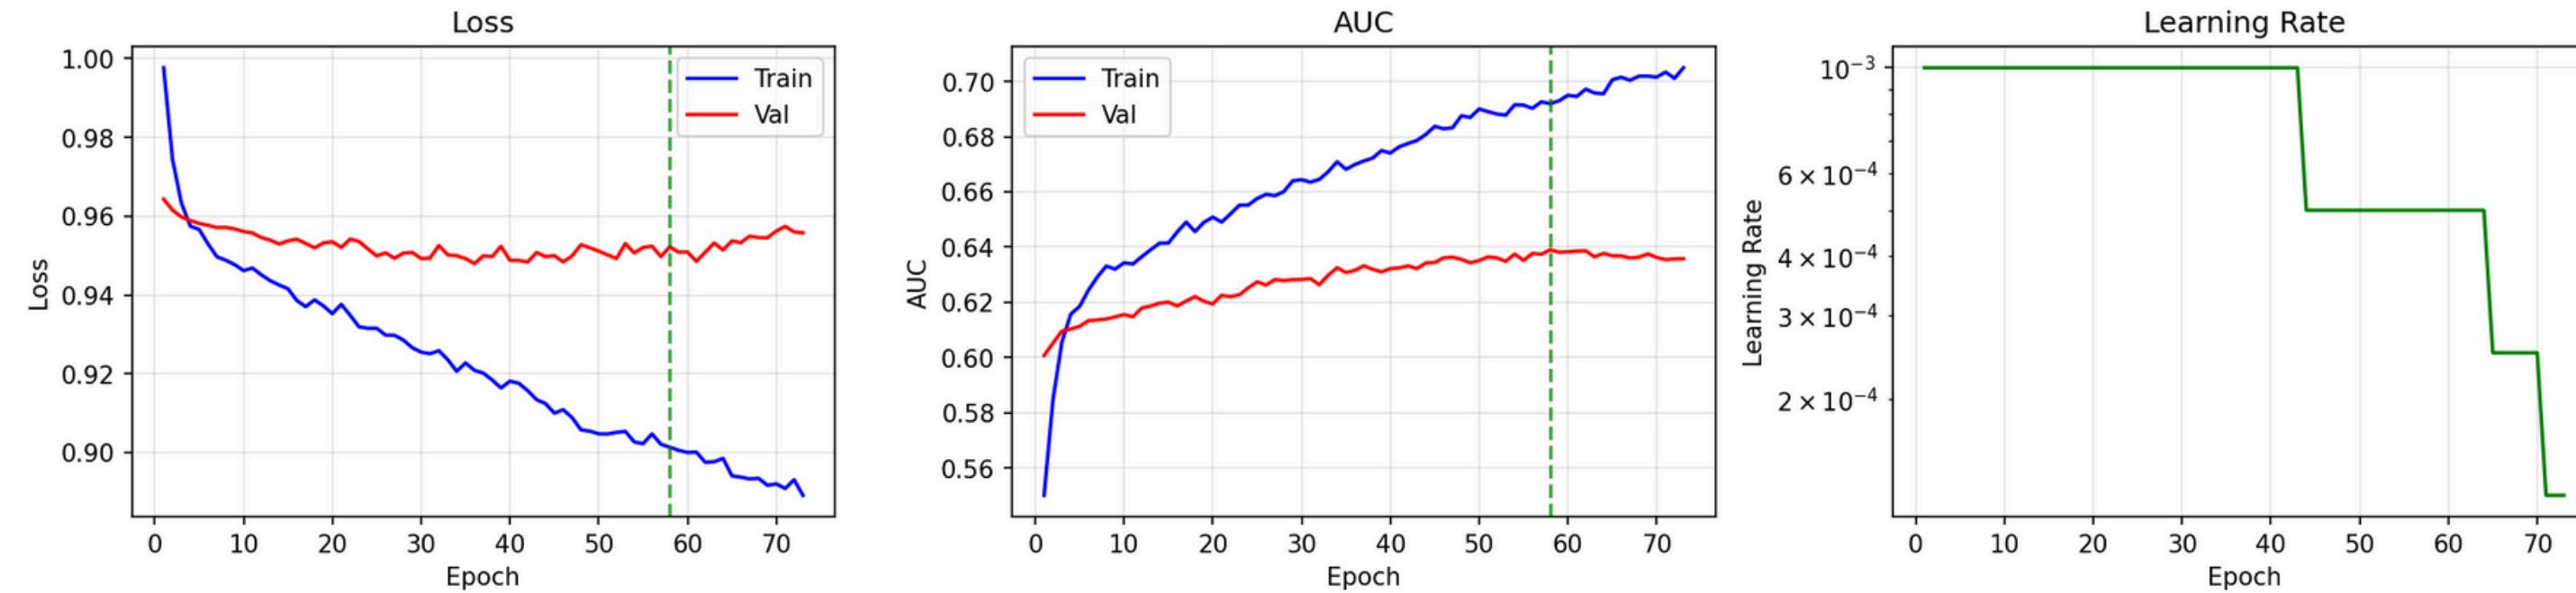


Model family	Examples in your project	Interpretability	Nonlinearity captured	Deployment complexity
Model family	Logistic Regression	High (coefficients, odds ratios)	Low (mostly linear unless manual interactions)	Low (fast, stable, easy to monitor)
Tree (single)	Decision Tree	Medium (rules, but unstable)	Medium (piecewise splits)	Low-Medium (simple but can overfit)
Tree ensemble (bagging)	Random Forest	Medium-Low (feature importance ok, less transparent)	High	Medium (bigger model, slower inference)
Gradient boosting	XGBoost / LightGBM	Medium (SHAP works well, but model is complex)	High	Medium (needs careful tuning + monitoring)
Deep learning (tabular)	Dense / Embedding / Wide&Deep (+ focal/weights)	Low (harder to explain; SHAP possible but heavier)	High	High (training/inference stack, reproducibility, drift)
Ensemble (stacking)	Stacking (meta-learner over base models)	Medium-Low (explain via meta-weights + SHAP per base)	High (combines diverse signals)	High (more components, OOF logic, versioning)

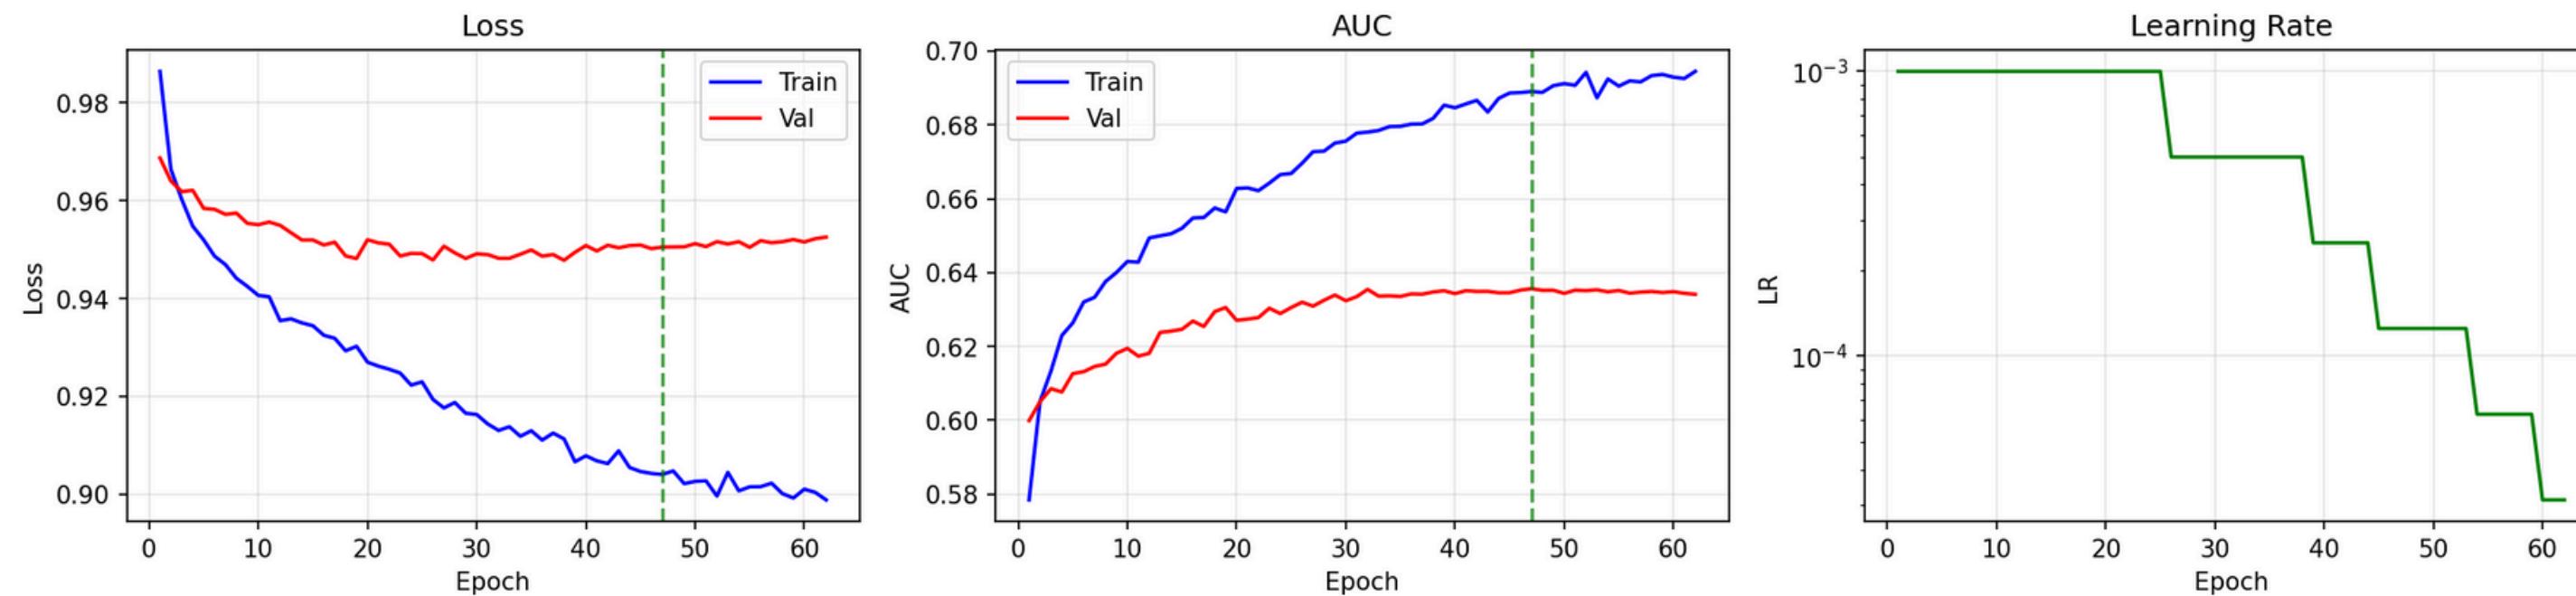
🏆 Model Performance Leaderboard (Test Set)

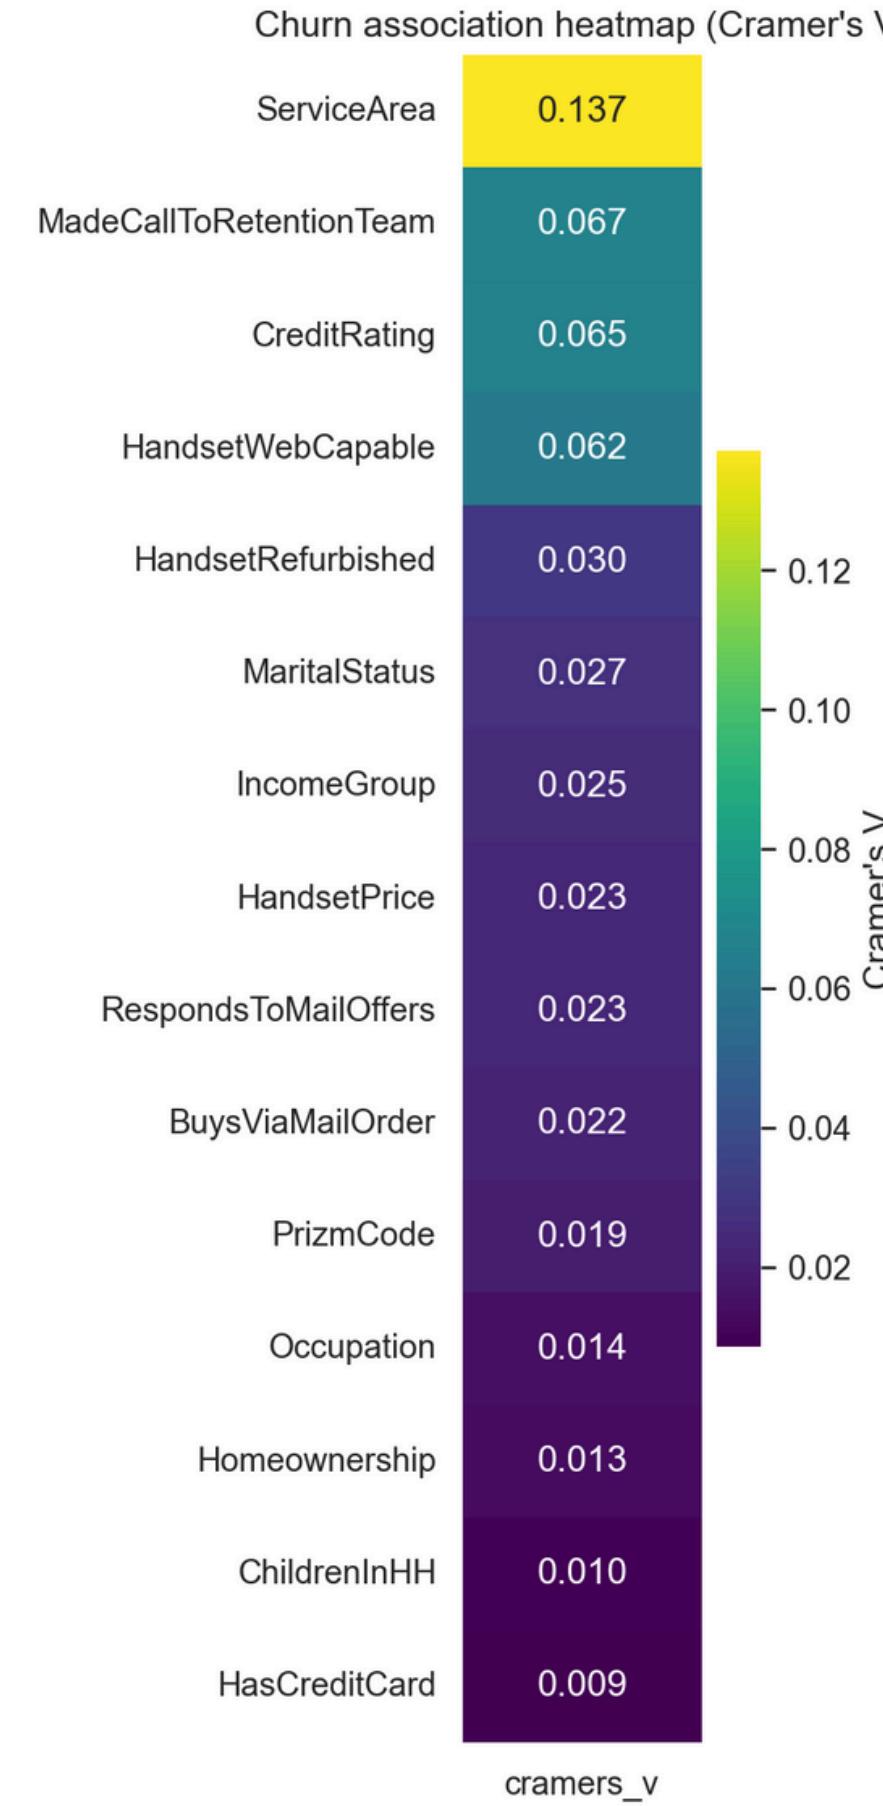
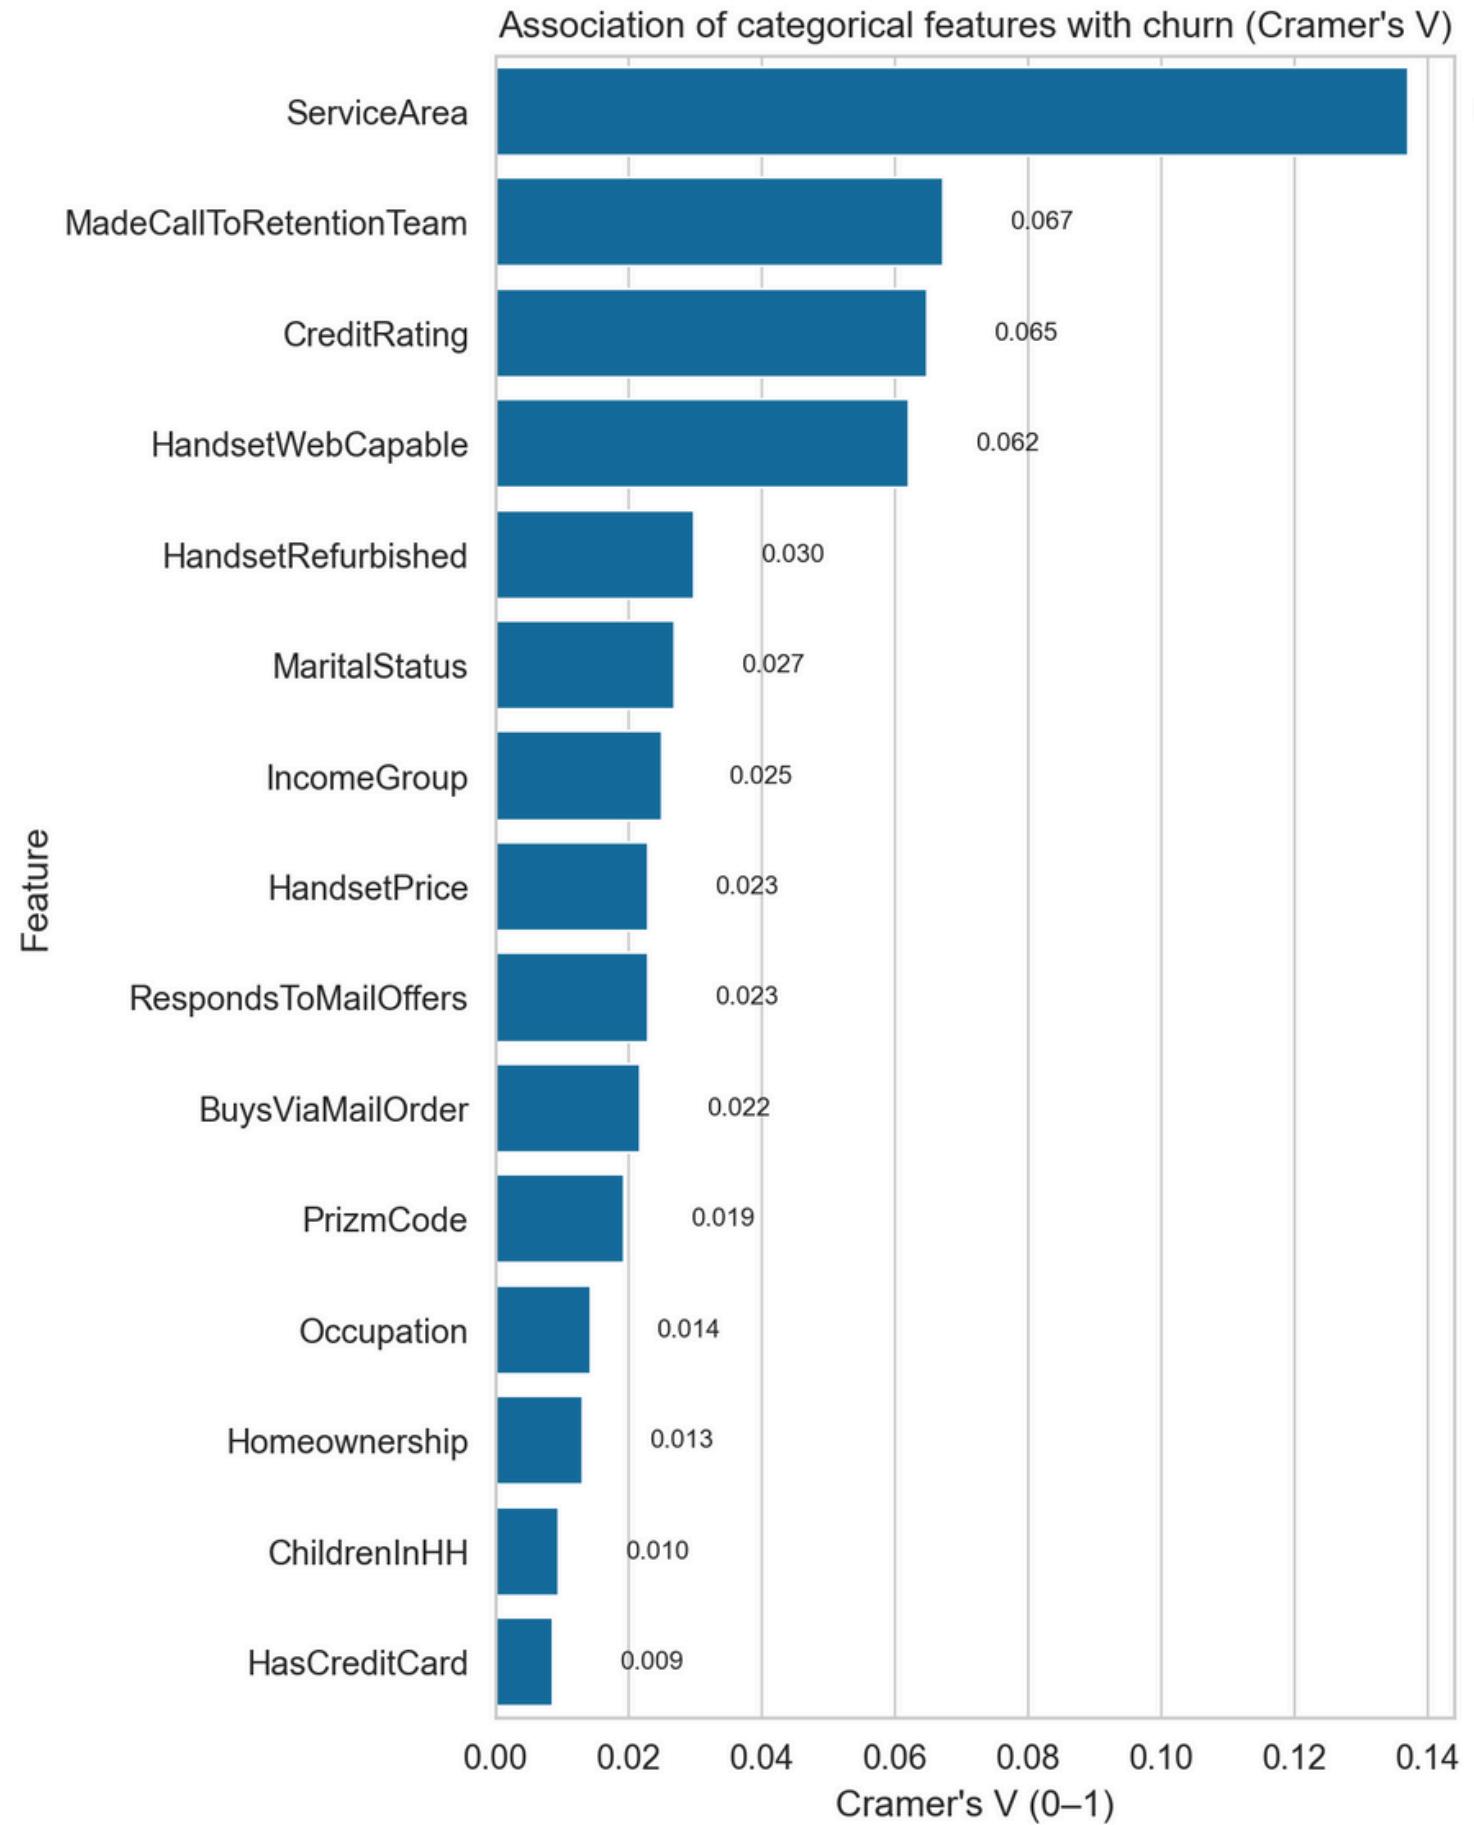
	Model	Family	Test AUC	Test PR-AUC	Test F1	Test Recall	Threshold
1	Stacking Ensemble	Ensemble	0.6773	0.4486	0.4965	0.6647	0.500
2	XGBoost (Optuna, thr=0.5)	GBM	0.6734	0.4484	0.1979	0.1205	0.500
3	XGBoost (Optuna, thr≈0.26)	GBM	0.6734	0.4484	0.4925	0.7331	0.258
4	XGBoost (Baseline)	GBM	0.6731	0.4440	0.2031	0.1242	0.500
5	LightGBM (Small Tuned)	GBM	0.6690	0.4448	0.4777	0.6017	0.500
6	Random Forest (Baseline)	Tree	0.6630	0.4306	0.4778	0.6053	0.500
7	NN - Wide & Deep (Focal)	NN	0.6584	0.4215	0.4852	0.6652	0.500
8	NN - Embedding (Focal)	NN	0.6575	0.4252	0.4782	0.6371	0.500
9	NN - Wide & Deep *	NN	0.6553	0.4259	0.4692	0.6094	0.500
10	NN - Dense (Baseline)	NN	0.6552	0.4192	0.4714	0.6176	0.500
11	NN - Embedding (Baseline)	NN	0.6549	0.4240	0.4737	0.6235	0.500
12	Decision Tree (Baseline)	Tree	0.6291	0.3794	0.4642	0.6588	0.500
13	Logistic Regression (Baseline)	Linear	0.6159	0.3774	0.4418	0.5677	0.500

Training Curves: embedding_deep

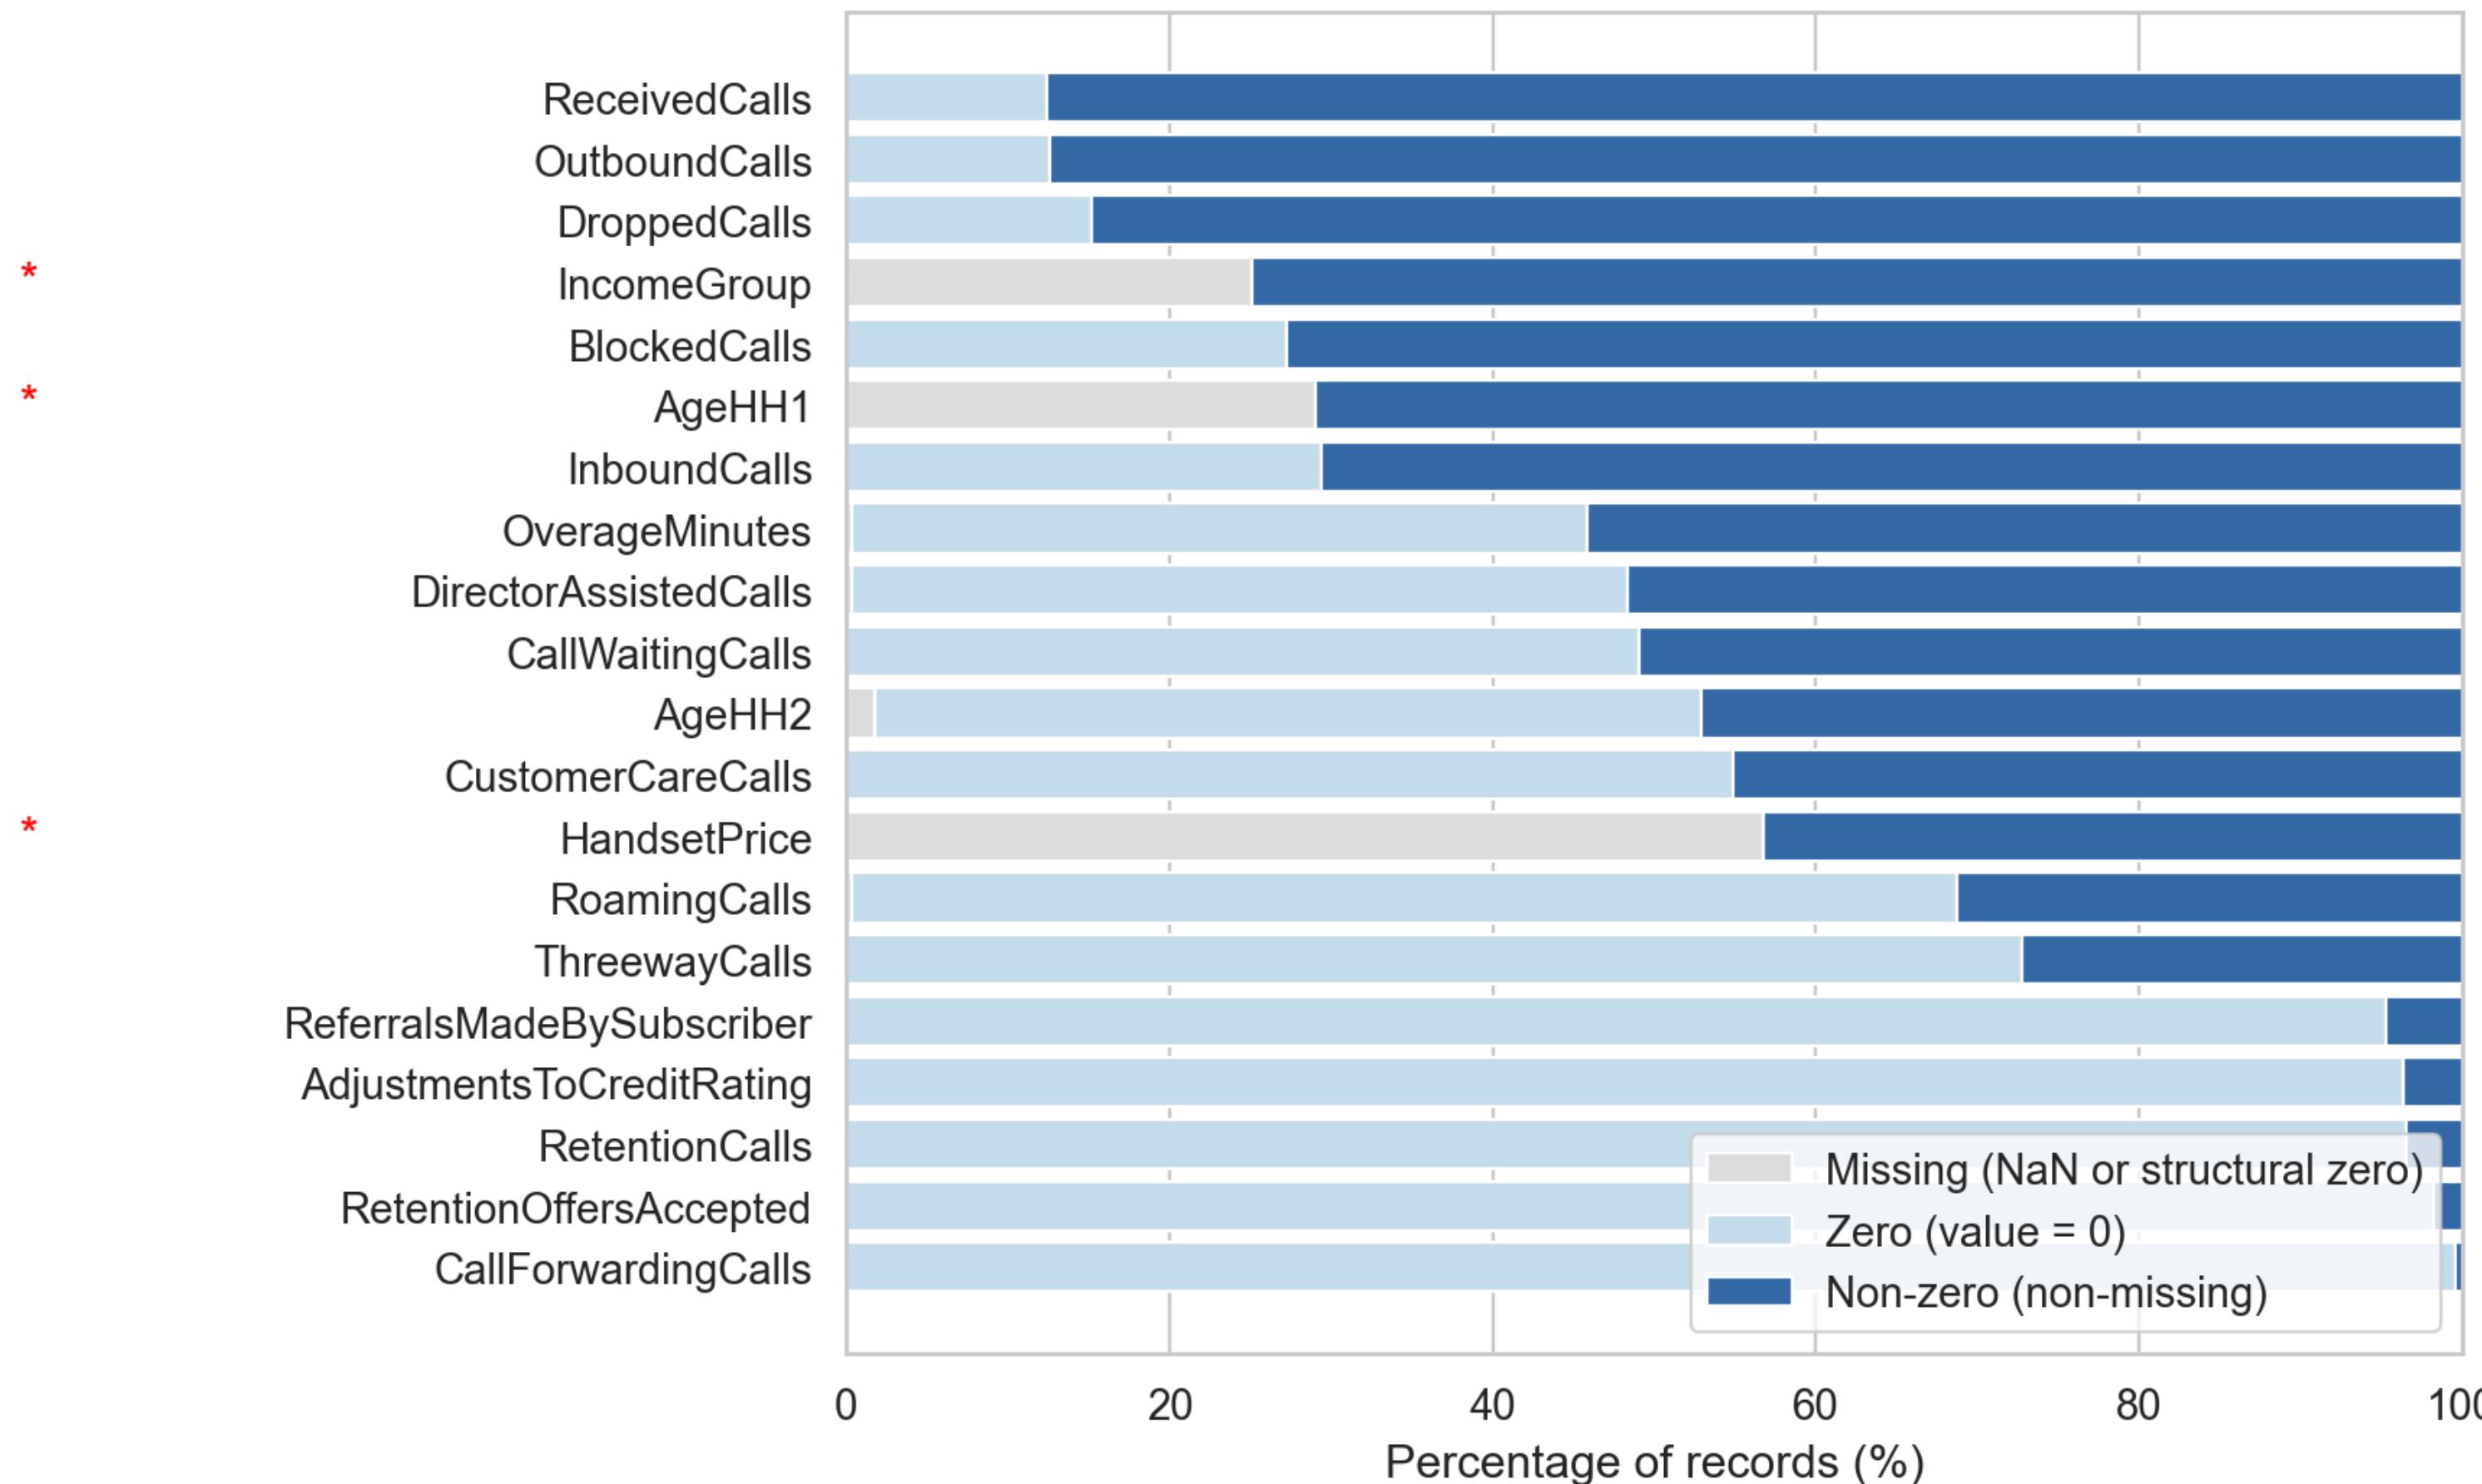


Training Curves: wide_and_deep





Missing and zero-value structure by feature (top sparsity)



* 0 is treated as missing for this feature

